

Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?

Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?
- “Time taken to perform a certain task T on system A is less than that taken to perform the same task T on system B.”
 - Note that this statement is specific to the task T. For some other task T', it may be possible that system B is the better performer.

Performance

<pre>int f = 1; for(int i=2; i<=n; i++) { f = f * i; } cout << f;</pre>	<pre>int f = 1; for(int i=2; i<=2*n; i++) { if (i <= n) f = f * i; } cout << f;</pre>
Program A	Program B

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- Which system displays higher performance?

Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

How many instructions are executed?

Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

- Number of static instructions = 8
- Number of dynamic instructions = 12

Performance

<pre>int f = 1; for(int i=2; i<=n; i++) { f = f * i; } cout << f;</pre>	<pre>int f = 1; for(int i=2; i<=2*n; i++) { if (i <= n) f = f * i; } cout << f;</pre>
Program A	Program B

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- System A displays higher performance because program A has fewer dynamic instructions (assuming every instruction takes equal amount of time)

SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?

SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?
- Work out the longest path in the combinational circuitry. Processor clock period should be greater than the time taken to execute this longest path. Let this clock period be t seconds. Processor frequency $f = 1/t$.
- Every instruction is executed in t s.
- Every instruction is executed in 1 cycle in our SimpleRISC processor.
 - Cycles per instruction, $CPI = 1$
 - Instructions per cycle, $IPC = 1$

Computing the Time a Program Takes

$$\begin{aligned}\tau &= \#seconds \\ &= \frac{\#seconds}{\#cycles} * \frac{\#cycles}{\#instructions} * (\#instructions) \\ &= \underbrace{\frac{\#seconds}{\#cycles}}_{1/f} + \underbrace{\frac{\#cycles}{\#instructions}}_{CPI} * (\#instructions) \\ &= \frac{CPI * \#insts}{f}\end{aligned}$$

- * **CPI** → **Cycles** per instruction
- * **f** → **frequency** (cycles per second)

The Performance Equation

$$P \propto \frac{IPC * f}{\#insts}$$

- * **IPC** → 1/CPI (Instructions per Cycle)
- * What are the **units** of performance ?
 - * **ANSWER** : arbitrary

Number of Instructions (#insts)

Static Instruction: The **binary** or executable of a **program**, contains a list of *static instructions*.

Dynamic Instruction: A **dynamic instruction** is a running instance of a static instruction, which is created by the **processor** when an instruction enters the pipeline.

- * Note that these are **dynamic** instructions
 - * **NOT** static instructions
- * A smart compiler can reduce the number of executed instructions

Number of Instructions(#insts) – 2

- * Dead code removal

- * Often **programmers** write **code** that does not determine the final output
- * This code is **redundant**
- * It can be identified and removed by the **compiler**


- * Function inlining

- * Very small **functions** have a lot of **overhead** → call, ret instructions, register spilling, and restoring
- * Paste the code of the **callee** in the code of the **caller** (known as **inlining**)

Computer Organisation and Architecture

**Smruti Ranjan Sarangi,
IIT Delhi**

Chapter 9 Principles of Pipelining

Also available as  Book

Computer Organisation and Architecture


Smruti Ranjan Sarangi



These slides are meant to be used along with the book: Computer Organisation and Architecture, Smruti Ranjan Sarangi, McGrawHill 2015
Visit: <http://www.cse.iitd.ernet.in/~srsarangi/archbooksoft.html>

Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Up till now

- * We have designed a **processor** that can execute all the **SimpleRisc** Instructions
- * We have look at two **styles** :
 - * With a **hardwired** control unit
 - * **Microprogrammed** control unit
 - * Microprogrammed data path
 - * Microassembly Language
 - * Microinstructions


Designing Efficient Processors

- * **Microprogrammed** processors are much slower than **hardwired** processors
- * Even **hardwired** processors
 - * Have a lot of **waste** !!!
 - * We have 5 stages.
 - * What is the IF stage doing, when the MA stage is active ?
 - * **ANSWER** : It is **idling**

The Notion of Pipelining

- * Let us go back to the **car assembly line**
 - * Is the **engine shop** idle, when the **paint shop** is painting a car ?
 - * **NO** : It is building the engine of another car
 - * When this engine goes to the body shop, it builds the engine of another car, and **so on**
- * **Insight** :
 - * **Multiple cars** are built at the same time.
 - * A car **proceeds** from one stage to the next

Pipelined Processors



- * The IF, ID, EX, MA, and RW stages process 5 instructions simultaneously
- * Each instruction proceeds from one stage to the next
- * This is known as **pipelining**

Advantages of Pipelining

- * We keep all parts of the data path, busy all the time
- * Let us assume that all the 5 stages do the same amount of **work**
 - * **Without** pipelining, every **T** seconds, an instruction completes its **execution**
 - * **With** pipelining, every **T/5** seconds, a new instruction completes its **execution**

Design of a Pipeline


* Splitting the Data Path

- * We **divide** the data path into 5 **parts** : IF, OF, EX, MA, and RW

* Timing

- * We insert **latches (registers)** between consecutive stages
- * 4 Latches → IF-OF, OF-EX, EX-MA, and MA-RW
- * At the **negative edge** of a clock, an **instruction** moves from one stage to the next

Pipelined Data Path with Latches




- * Add a latch between subsequent stages.
 - * Triggered by a negative clock edge

The Instruction Packet

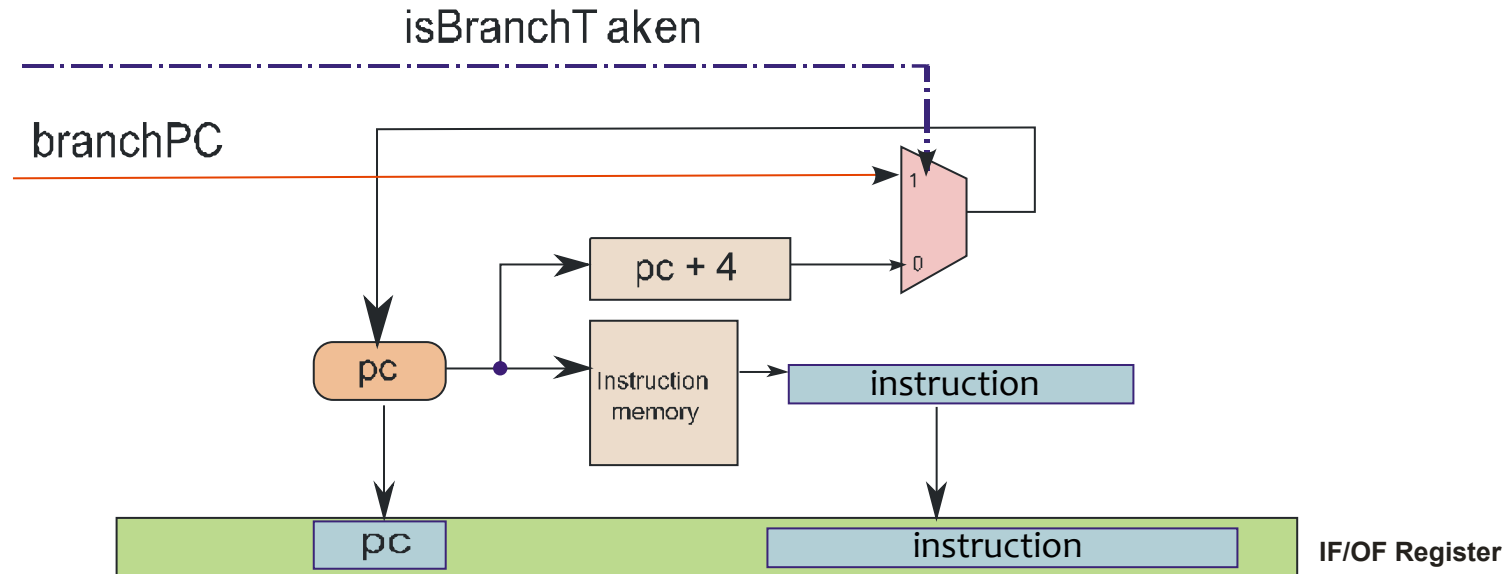
- * What travels **between stages** ?
 - * **ANSWER** : the instruction packet
- * Instruction Packet
 - * **Instruction contents**
 - * Program counter
 - * **All intermediate results**
 - * Control signals
- * Every instruction moves with its entire state, no interference between instructions

Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions

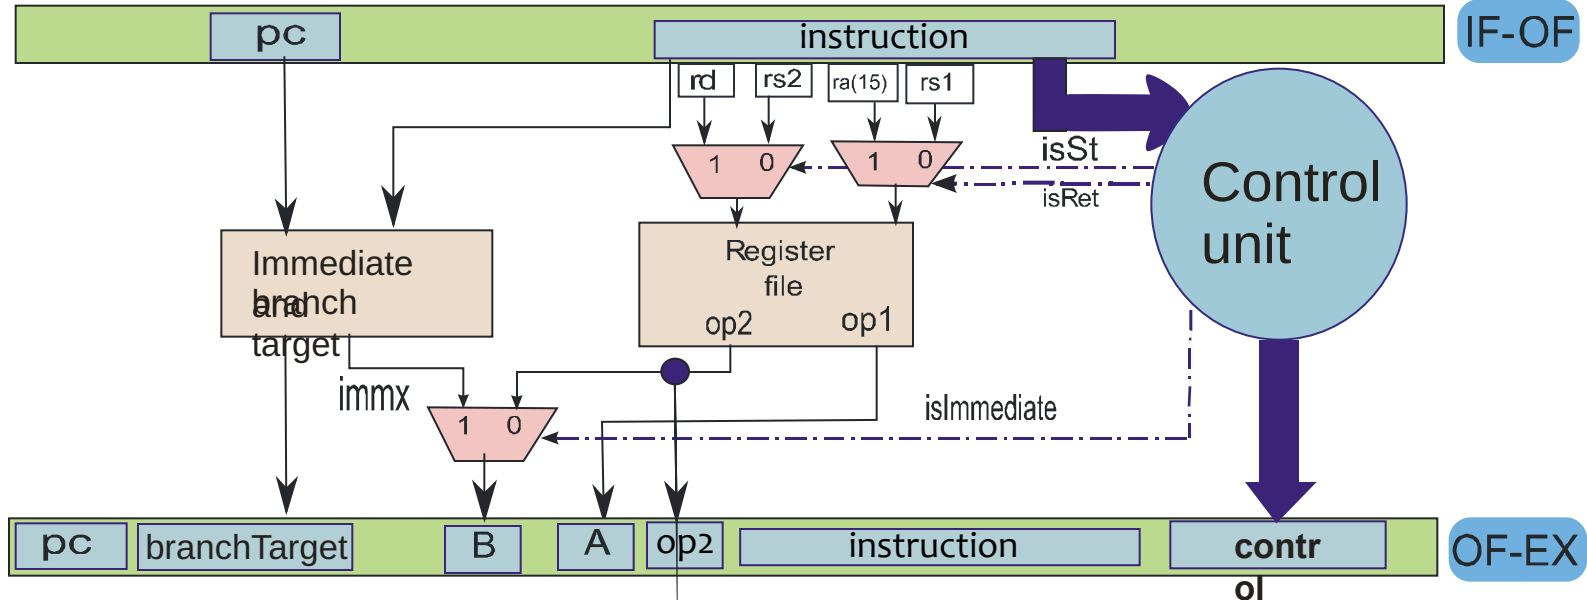


IF Stage



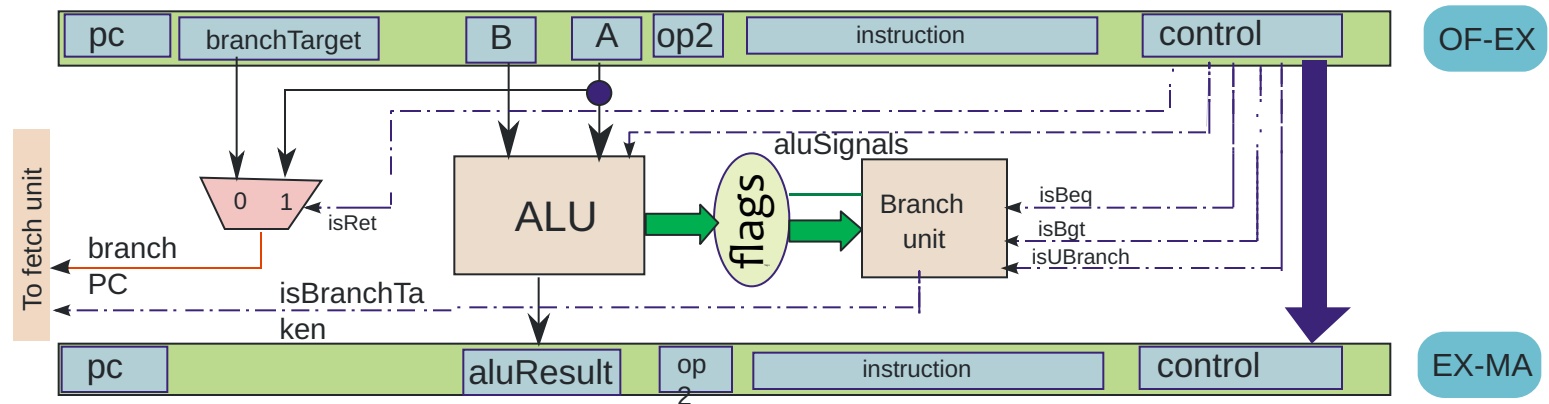
- * Instruction contents saved in the **instruction** field

OF Stage



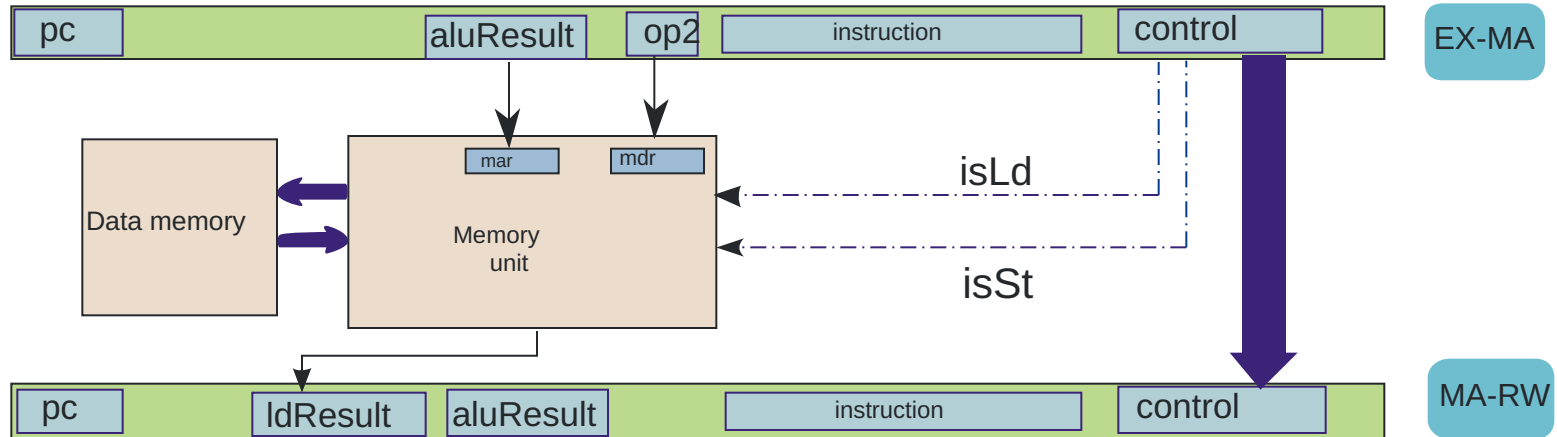
* **A, B** → ALU Operands, **op2** (store operand),
control (set of all control signals)

EX Stage



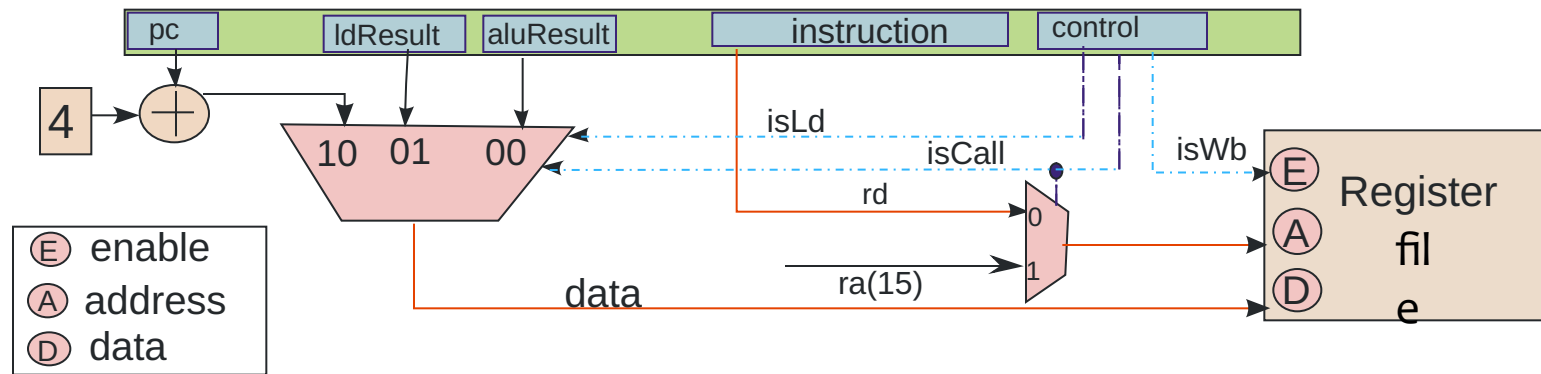
- * **aluResult** → result of the ALU Operation
- * **op₂, control, pc, instruction** (passed from OF-EX)

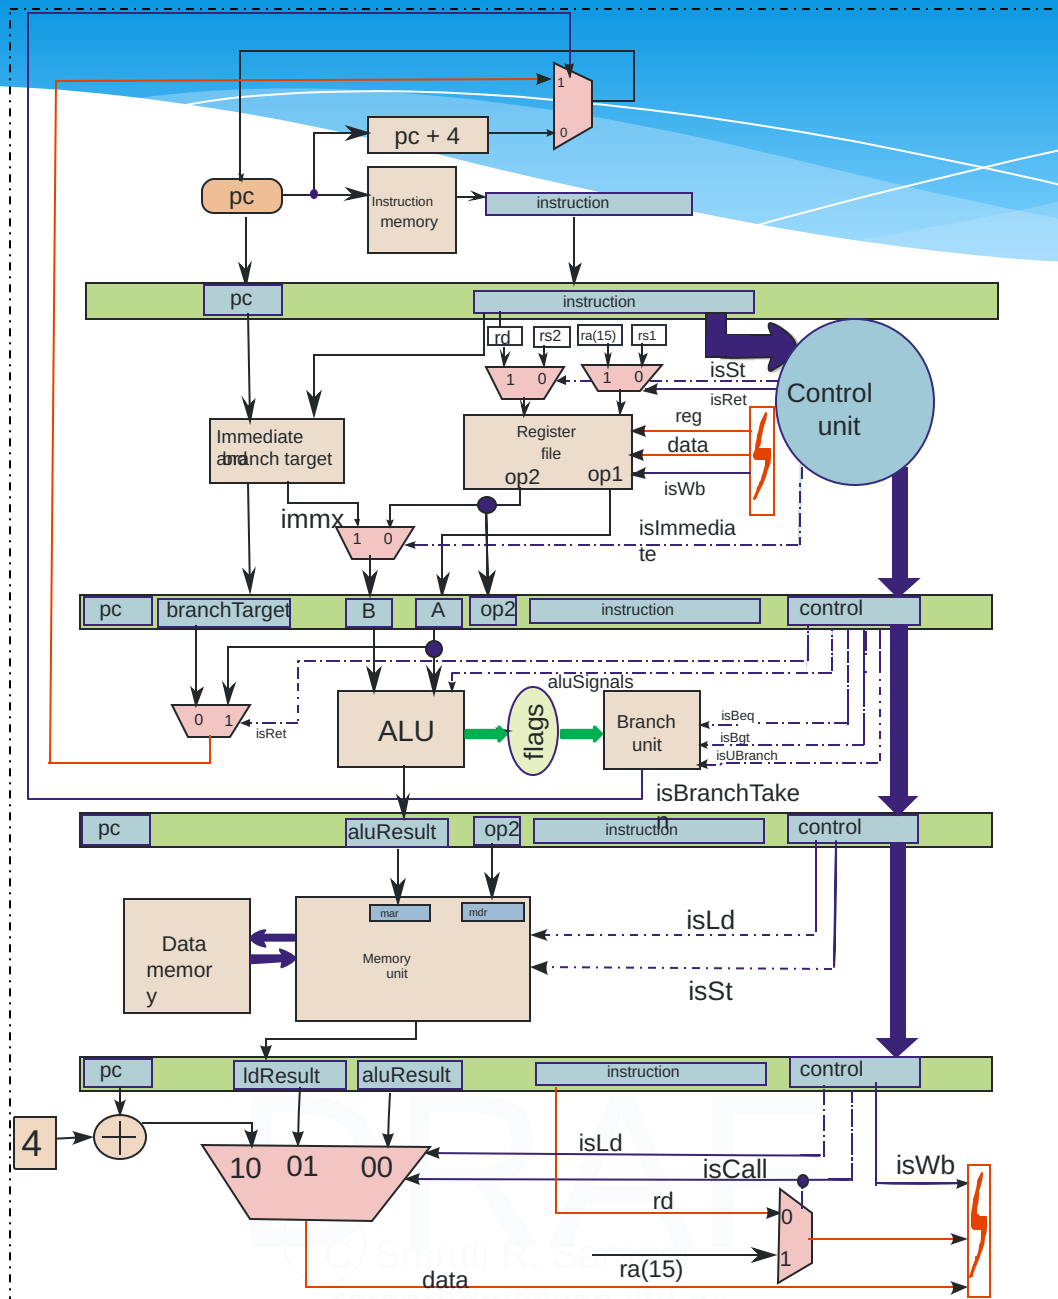
MA Stage



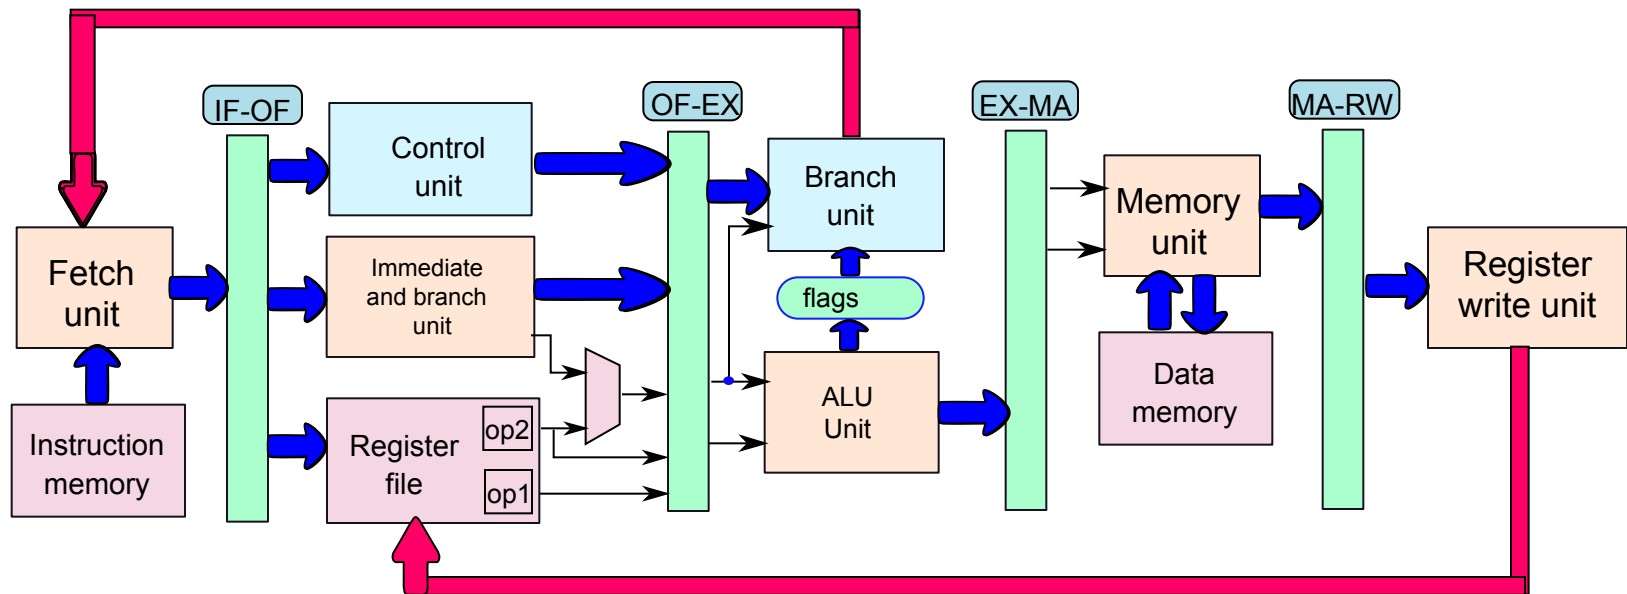
- * **IdResult** → result of the load operation
- * **aluResult, control, pc, instruction** (passed from EX-MA)

RW Stage






Abridged Diagram



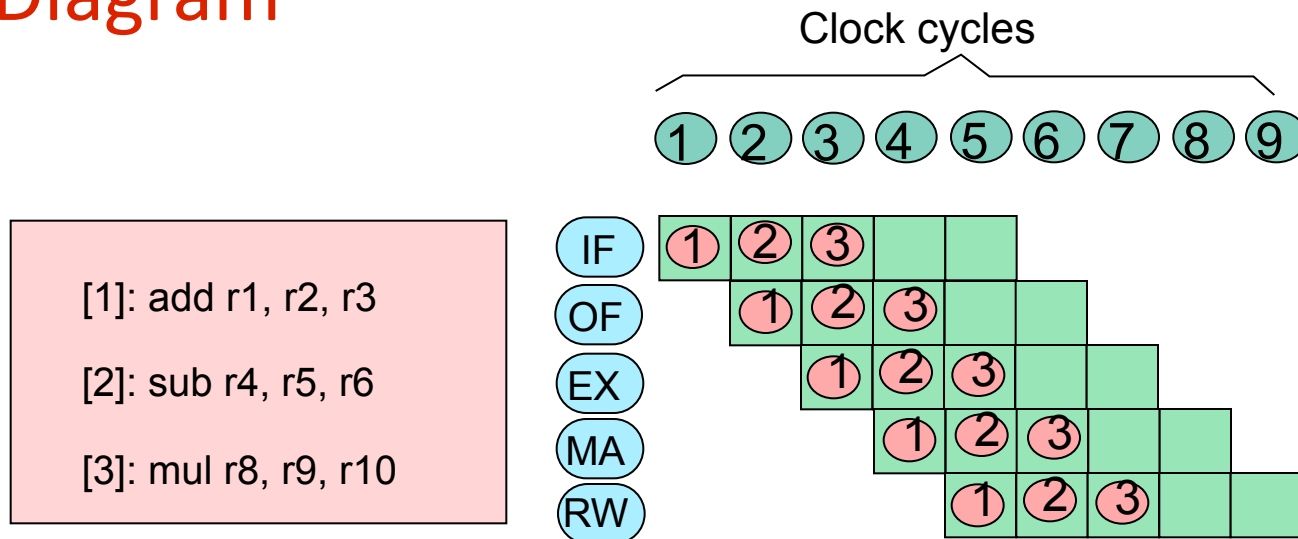
Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Pipeline Hazards

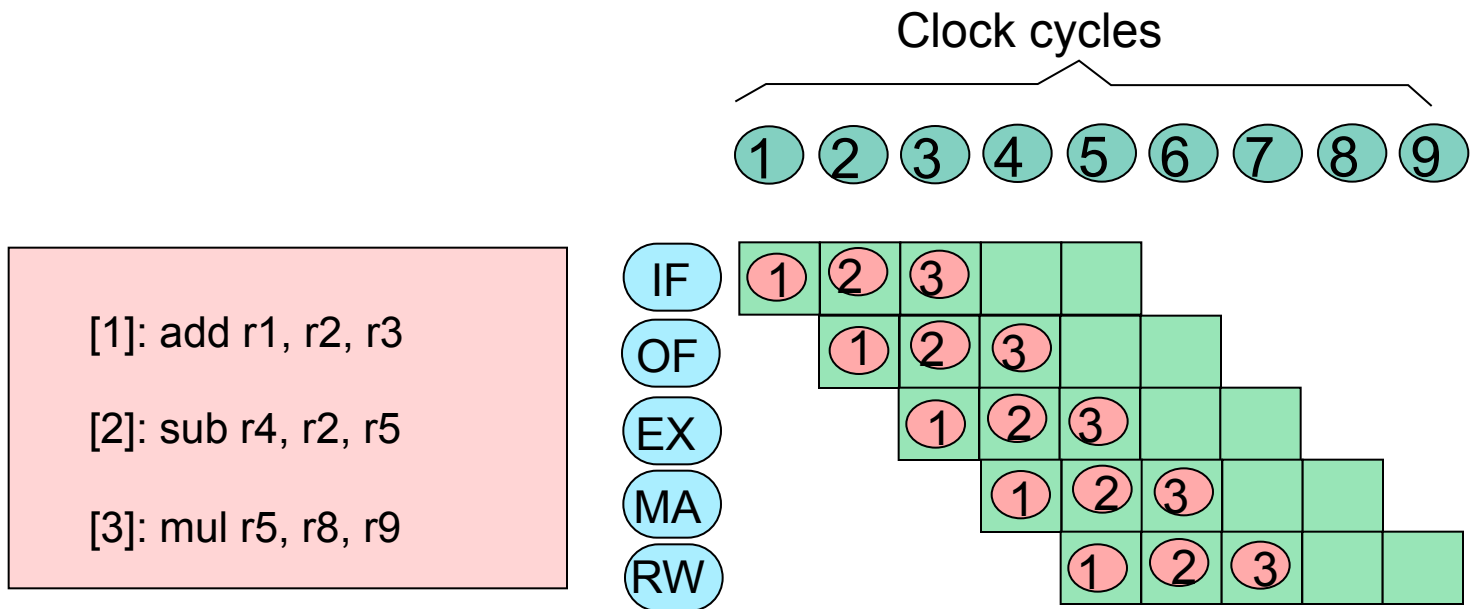
- * Now, let us consider **correctness**
- * Let us introduce a new tool → **Pipeline Diagram**



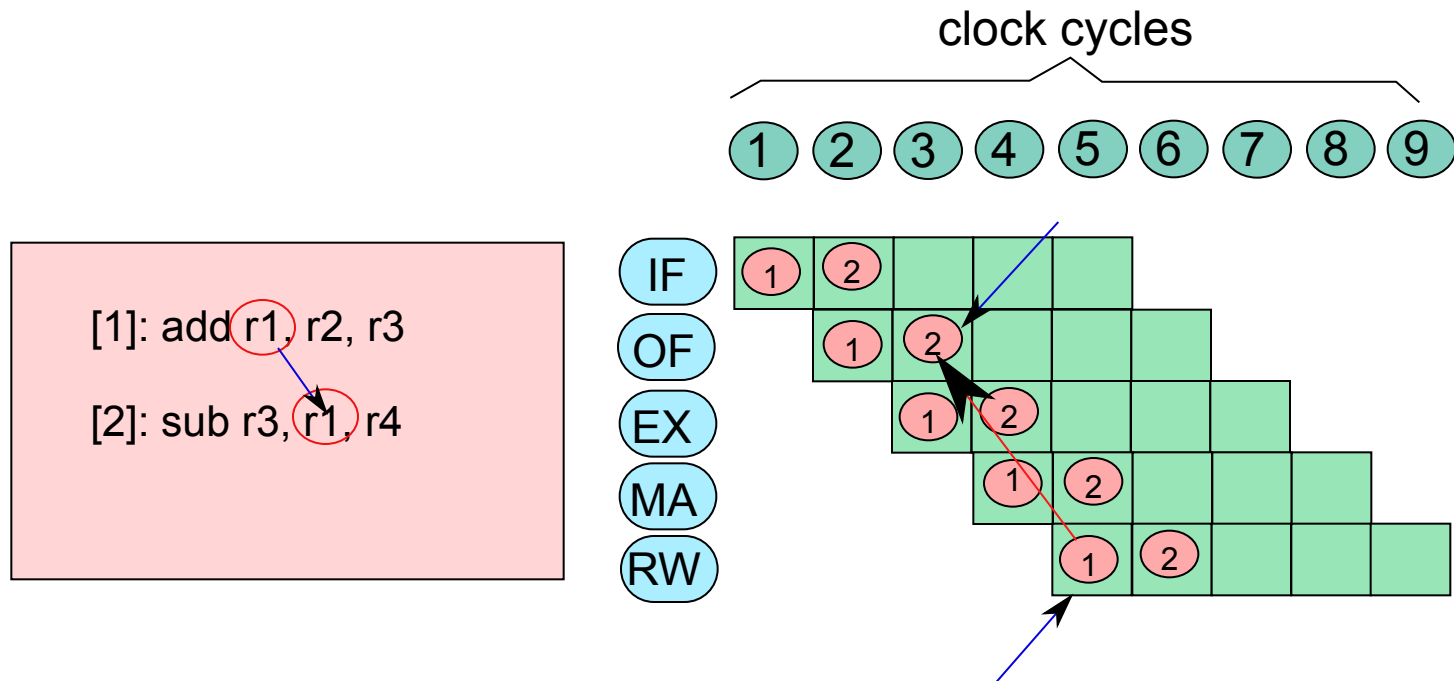
Rules for Constructing a Pipeline Diagram

- * It has 5 **rows**
 - * One per each stage
 - * The rows are **named** : IF, OF, EX, MA, and RW
- * Each **column** represents a clock cycle
- * Each **cell** represents the execution of an instruction in a stage
 - * It is **annotated** with the name(label) of the instruction
- * Instructions proceed from one stage to the next across clock cycles

Example



Data Hazards



* Instruction 2 will read incorrect values !!!

Data Hazard

Definition: A **hazard** is defined as the possibility of erroneous execution of an instruction in a pipeline. A data hazard represents the possibility of erroneous execution because of the unavailability of data, or the availability of **incorrect** data.

- * This situation represents a **data hazard**
- * In specific,
 - * it is a **RAW** (read after write) hazard
- * The earliest we can dispatch instruction 2, is cycle 5

Other Types of Data Hazards

- * Our pipeline is in-order

Definition: In an in-order pipeline (such as ours), a preceding instruction is always ahead of a succeeding instruction in the pipeline. Modern processors however use out-of-order pipelines that break this rule. It is possible for later instructions to execute before earlier instructions.

- * We will only have RAW hazards in our pipeline.
- * Out-of-order pipelines can have WAR and WAW hazards

WAW Hazards

```
[1]: add r1, r2, r3  
[2]: sub r1, r4, r3
```

- * Instruction [2] cannot write the value of r1, before instruction [1] writes to it, will lead to a **WAW** hazard

WAR Hazards

```
[1]: add r1, r2, r3  
[2]: add r2, r5, r6
```

- * **Instruction** [2] cannot **write** the value of r2, before **instruction** [1] reads it → will lead to a **WAR** hazard

[1] add r1, r2, r3
[2] ld r4, 8[r1]
[3] sub r5, r5, r2
[4] st r4, 12[r2]
[5] sub r6, r4, r3
[6] add r7, r4, r3

Q. How many dependencies?

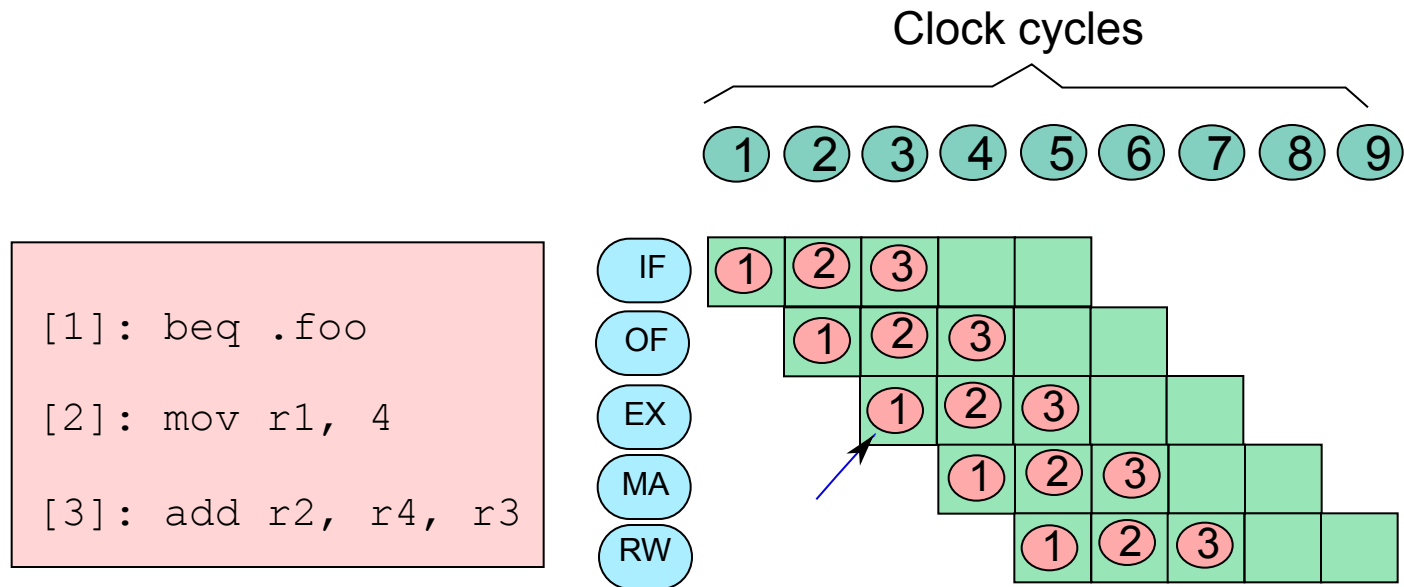
Q. How many dependencies require conflict resolution?

Control Hazards

```
[1]: beq .foo  
[2]: mov r1, 4  
[3]: add r2, r4, r3  
...  
...  
.foo:  
[100]: add r4, r1, r2
```

- * If the branch is **taken**, instructions [2] and [3], might get fetched, incorrectly

Control Hazard – Pipeline Diagram



- * The two **instructions** fetched immediately after a branch instruction might have been fetched **incorrectly**.

```
[1] mov r1, 0
[2] mov r2, 0
[3] cmp r1, r2
[4] beq .foo
[5] sub r6, r4, r3
[6] add r7, r4, r3
[7] mul r8, r4, r3
...
...
.foo
[100] add r5, r4, r1
```

Control Hazards

- * The two **instructions** fetched immediately after a branch instruction might have been fetched **incorrectly**.
- * These instructions are said to be on the **wrong path**
- * A **control hazard** represents the **possibility** of **erroneous** execution in a **pipeline** because instructions in the **wrong path** of a **branch** can possibly get **executed** and save their results in memory, or in the register file

Structural Hazards

- * A **structural hazard** may occur when two instructions have a conflict on the same set of resources in a cycle
- * Example :
 - * Assume that we have an add instruction that can read one operand from memory
 - * `add r1, r2, 10[r3]`

Structural Hazards - II

```
[1]: st r4, 20[r5]  
[2]: sub r8, r9, r10  
[3]: add r1, r2, 10[r3]
```


- * This code will have a structural hazard
 - * [3] tries to read 10[r3] (MA unit) in cycle 4
 - * [1] tries to write to 20[r5] (MA unit) in cycle 4
- * Does not happen in our pipeline

Solutions in Software

- * Data hazards

- * Insert **nop** instructions, reorder **code**

```
[1]: add r1, r2, r3  
[2]: sub r3, r1, r4
```




```
[1]: add r1, r2, r3  
[2]: nop  
[3]: nop  
[4]: nop  
[5]: sub r3, r1, r4
```


Code Reordering

```
add r1, r2, r3  
add r4, r1, 3  
add r8, r5, r6  
add r9, r8, r5  
add r10, r11, r12  
add r13, r10, 2
```

Code Reordering

```
add r1, r2, r3
add r4, r1, 3
add r8, r5, r6
add r9, r8, r5
add r10, r11, r12
add r13, r10, 2
```




```
add r1, r2, r3
add r8, r5, r6
add r10, r11, r12
nop
add r4, r1, 3
add r9, r8, r5
add r13, r10, 2
```

Control Hazards

- * **Trivial Solution** : Add two nop instructions after every branch
- * **Better solution** :
 - * Assume that the two instructions fetched after a branch are **valid** instructions
 - * These instructions are said to be in the **delay slots**
 - * Such a branch is known as a **delayed branch**

Example with 2 Delay Slots

```
add r1, r2, r3
add r4, r5, r6
b .foo
add r8, r9, r10
```




```
b .foo
add r1, r2, r3
add r4, r5, r6
add r8, r9, r10
```

- * The **compiler** transfers **instructions** before the branch to the **delay slots**.
- * If it cannot find 2 valid **instructions**, it inserts **nops**.

Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Why interlocks ?

- * We cannot always **trust** the **compiler** to do a good job, or even introduce **nop** instructions correctly.
- * **Compilers** now need to be tailored to specific hardware.
- * We should ideally not expose the details of the **pipeline** to the **compiler** (might be confidential also)
- * Hardware mechanism to enforce correctness → **interlock**

Two kinds of Interlocks

- * Data-Lock

- * Do not allow a consumer instruction to move beyond the OF stage till it has read the correct values. Implication : Stall the IF and OF stages.

- * Branch-Lock

- * We never execute instructions in the wrong path.
- * The hardware needs to ensure both these conditions.

Comparison between Software and Hardware

Attribute	Software	Hardware(withinterlocks)
Portability	Limited to a specific processor	Programs can be run on any processor irrespective of the nature of the pipeline
Branches	Possible to have no performance penalty, by using delay slots	Need to stall the pipeline for 2 cycles in our design
RAW hazards	Possible to eliminate them through code scheduling	Need to stall the pipeline
Performance	Highly dependent on the nature of the program	The basic version of a pipeline with interlocks is expected to be slower than the version that relies on software

Conceptual Look at Pipeline with Interlocks

```
[1]: add r1, r2, r3  
[2]: sub r4, r1, r2
```

- * We have a **RAW hazard**
- * We need to **stall**, instruction [2] at the OF stage for 3 cycles.
- * We need to keep sending **nop** instructions to the **EX stage** during these 3 cycles


Example

 bubble


[1]: add r1, r2, r3

[2]: sub r4, r1, r2

Clock cycles




Example



[1]: add r1, r2, r3
[2]:
[3]: sub r4, r1, r2

IF
OF
EX
MA
RW

Clock cycles




A Pipeline Bubble

- * A **pipeline bubble** is inserted into a stage, when the **previous stage** needs to be stalled
- * It is a **nop** instruction
- * To insert a **bubble**
 - * Create a **nop** instruction packet
 - * OR, Mark a designated **bubble bit** to 1

Bubbles in the Case of a Branch Instruction



```
[1]: beq. foo
[2]: add r1, r2, r3
[3]: sub r4, r5, r6
....
....
.foo:
[4]: add r8, r9, r10
```



Control Hazards and Bubbles

- * We know that an instruction is a **branch** in the OF stage
- * When it reaches the EX stage and the branch is taken, let us **convert** the instructions in the IF, and OF stages to **bubbles**
- * Ensures the **branch-lock** condition

IF

IF-OF Latch

Inst_OF

OF

OF-EX Latch

Inst_EX

EX

EX-MA Latch

Inst_MA

MA

MA-RW Latch

Inst_RW

RW

Ensuring the Data-Lock Condition

- * When an **instruction** reaches the **OF** stage, check if it has a **conflict** with any of the instructions in the **EX, MA, and RW** stages
- * If there is **no conflict**, **nothing** needs to be done
- * Otherwise, **stall** the **pipeline** (IF and OF stages only)

Algorithm 5: Algorithm to detect conflicts between instructions

Data: instructions, [A], and [B]

Result: conflict exists (**true**), no conflict (**false**)

```
if [A].opcode  $\in$  (nop,b,beq,bgt,call) then
    /* Does not read from any register */
    return false
end
if [B].opcode  $\in$  (nop, cmp, st, b, beq, bgt, ret) then
    /* Does not write to any register */
    return false
end
/* Set the sources */
src1  $\leftarrow$  [A].rs1
src2  $\leftarrow$  [A].rs2
if [A].opcode = st then
    src2  $\leftarrow$  [A].rd
end
if [A].opcode = ret then
    src1  $\leftarrow$  ra
end
hasSrc1  $\leftarrow$  true
if ([A]  $\in$  (not, mov)) hasSrc1  $\leftarrow$  false
```

```

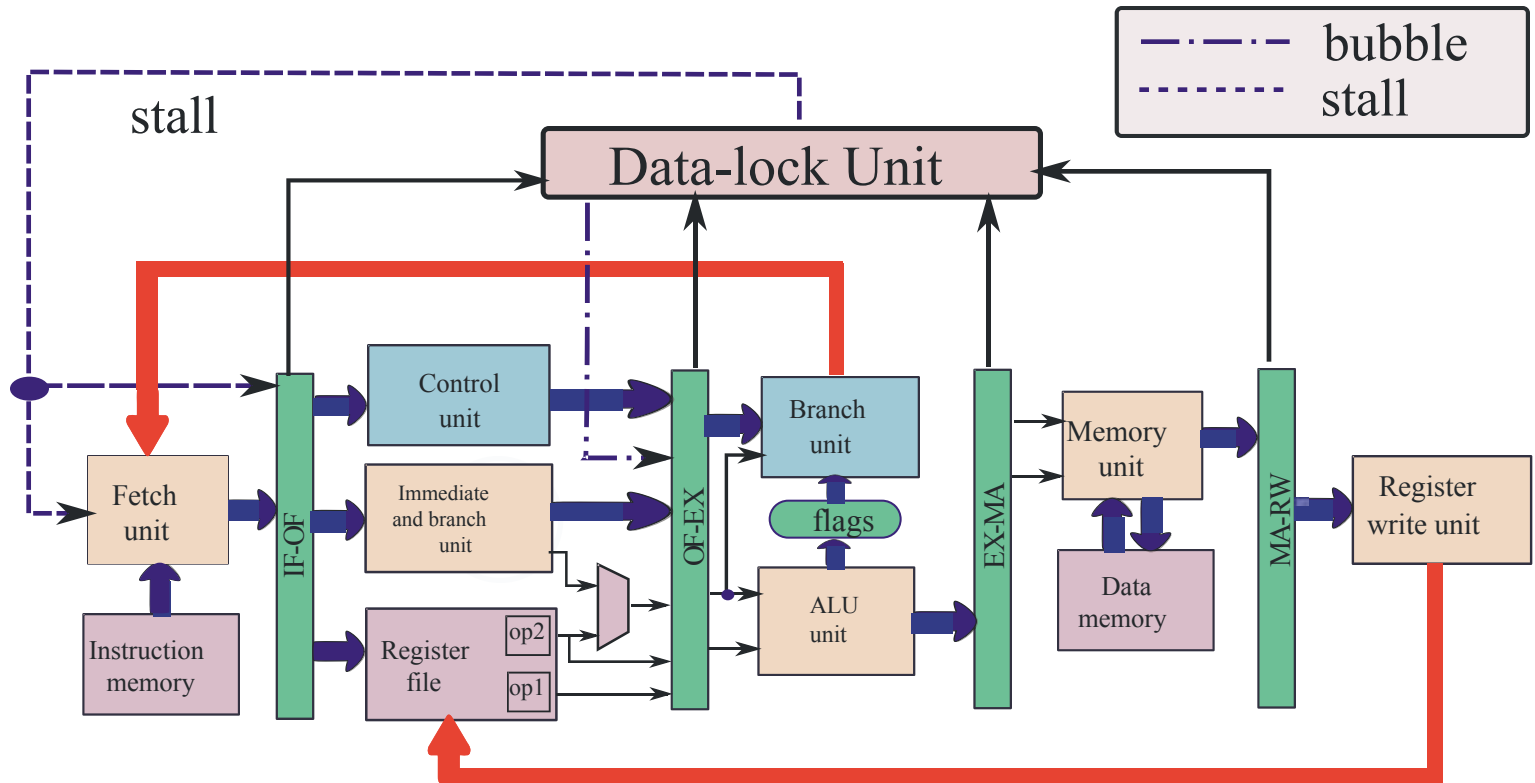
dest ← [B].rd
if [B].opcode = call then
    dest ← ra
end
/* Check the second operand to see if it is a register */
hasSrc2 ← true
if [A].opcode ≠ (st) then
    if [A].I = 1 then
        hasSrc2 ← false
    end
end
/* Detect conflicts */
if (hasSrc1 = true) and (src1 = dest) then
    return true
end
else if (hasSrc2 = true) and (src2 = dest) then
    return true
end
return false

```

How to Stall a Pipeline ?

- * Disable the **write functionality** of :
 - * The **IF-OF** register
 - * and the **Program Counter (PC)**
- * To insert a **bubble**
 - * Write a **bubble** (**nop** instruction) into the **OF-EX** register

Data Path with Interlocks (Data-Lock)



Ensuring the Branch-Lock Condition

- * Option 1 :

- * Use **delay slots** (**interlocks** not required)

- * Option 2 :

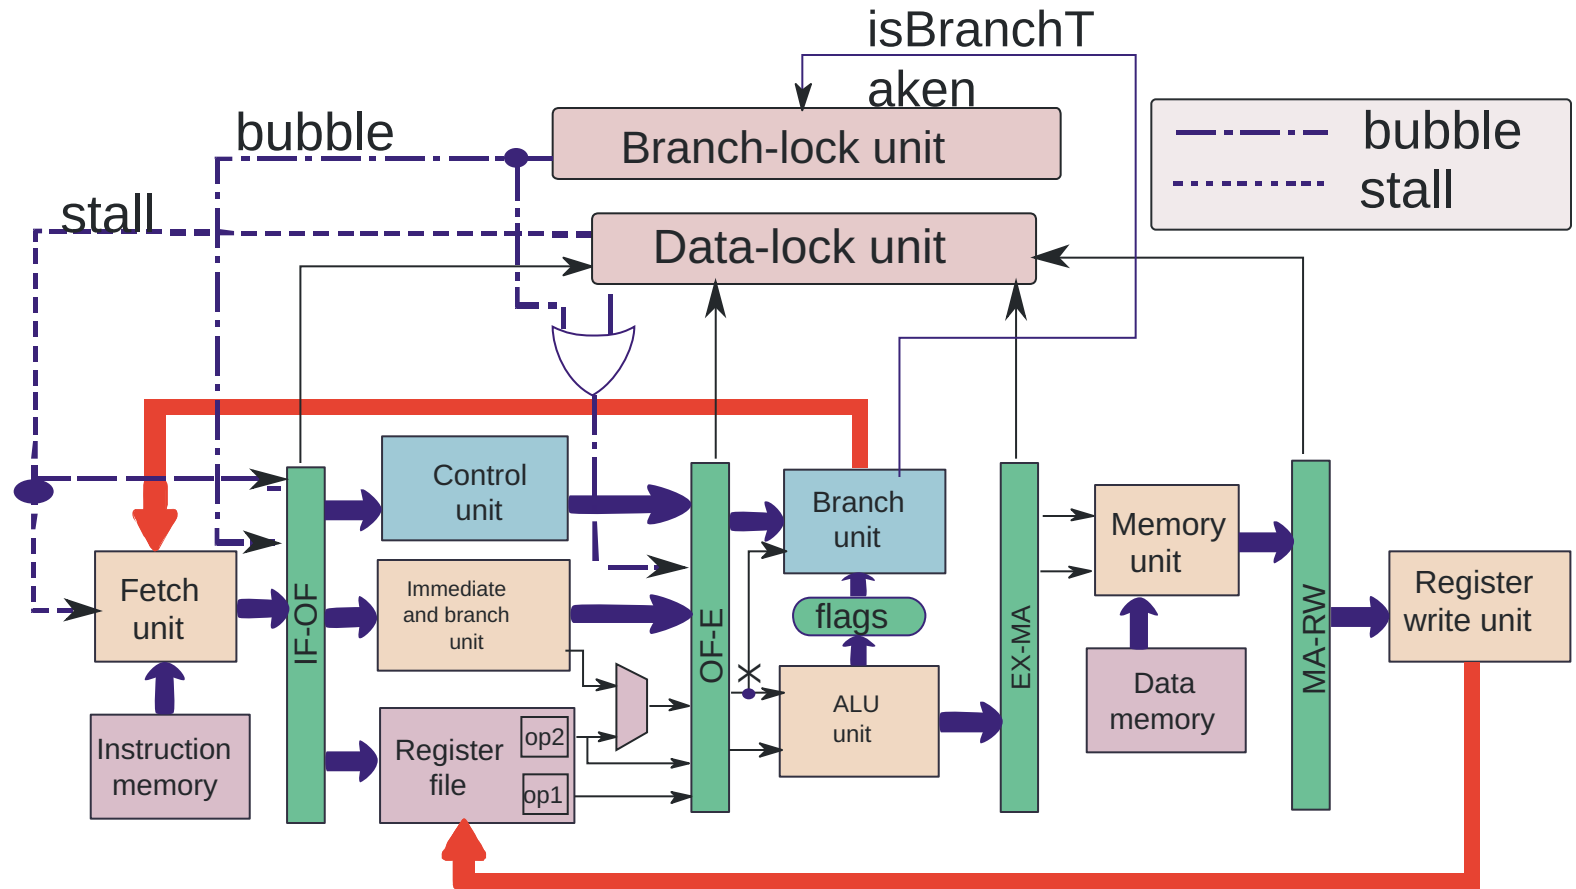
- * **Convert** the instructions in the IF, and OF stages, to **bubbles** once a **branch instruction** reaches the **EX** stage.
 - * Start fetching from the **next PC (not taken)** or the **branch target (taken)**

Ensuring the Branch-Lock Condition - II

* Option 3


- * If the **branch instruction** in the **EX** stage is **taken**, then **invalidate** the instructions in the IF and OF stages. Start **fetching** from the **branch target**.
- * Otherwise, do not take **any special action**
- * **This method is also called predict not-taken (we shall use this method because it is more efficient than option 2)**

Data Path with Interlocks




Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Relook at the Pipeline Diagram

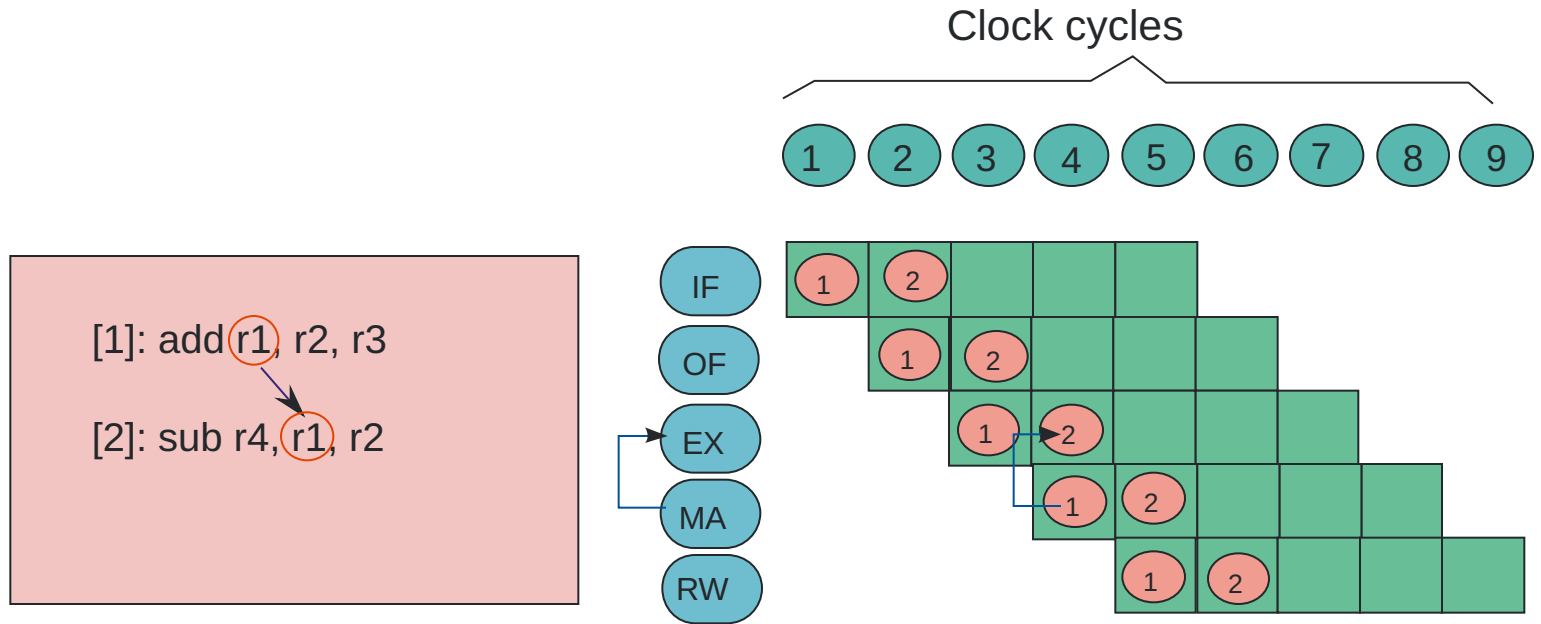


- * (a) → with bubbles
- * (b) → no bubbles (may lead to wrong results)

Crucial Insight (Figure (b))

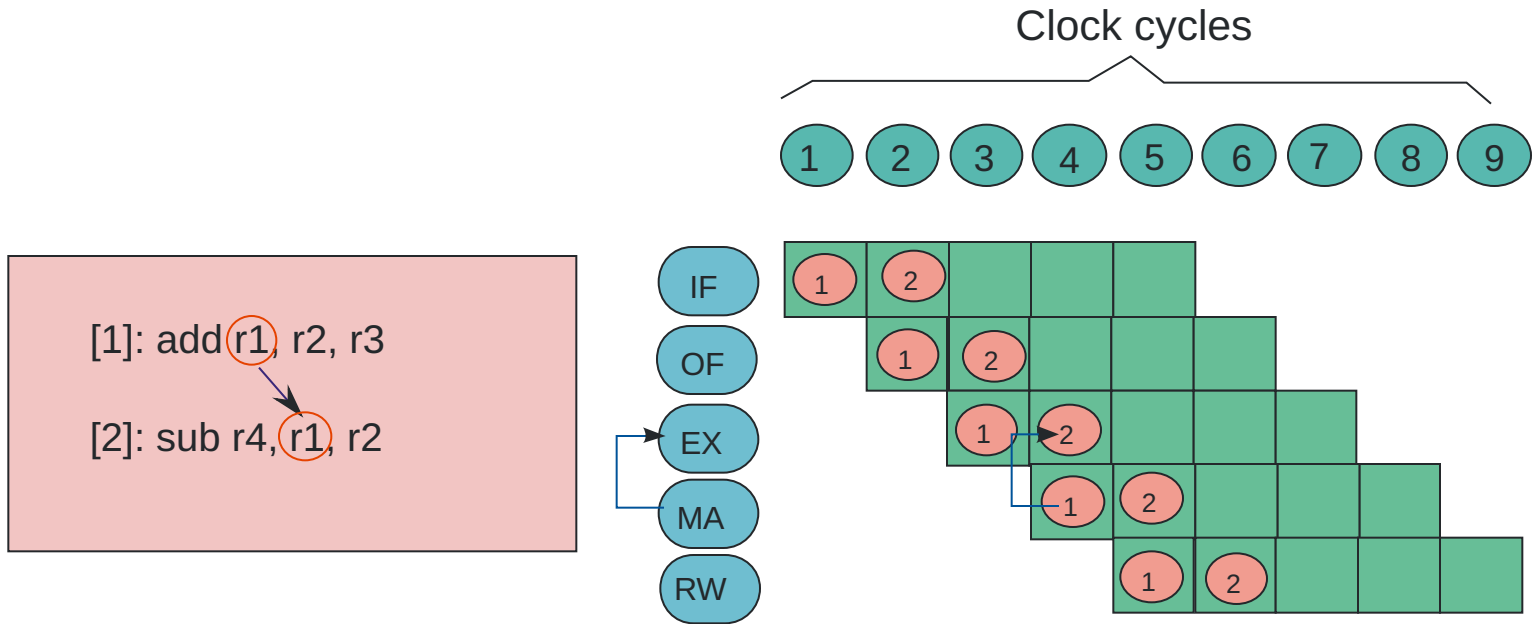
- * When does instruction 2 need the value of r1 ?
 - * **ANSWER** : Cycle 3, OF Stage (**wrong!!!**)
 - * **CORRECT ANSWER** : Cycle 4, EX Stage
- * When does instruction 1 produce the value of r1 ?
 - * **ANSWER** : Cycle 5, RW Stage (**wrong!!!**)
 - * **CORRECT ANSWER** : End of Cycle 3, EX Stage

Forwarding



- * If the **correct value** is already there in another stage, we can **forward** it.

Forwarding from MA to EX



- * Forwarding in cycle 4 from instruction [1] to [2]

Different Forwarding Paths

- * We need to add a multitude of forwarding paths
- * Rules for creating forwarding paths
 - * Add a path from a later stage to an earlier stage
 - * Try to add a forwarding path as late as possible. For, example, we avoid the $EX \rightarrow OF$ forwarding path, since we have the $MA \rightarrow EX$ forwarding path
 - * The IF stage is not a part of any forwarding path.

Forwarding Path

- * 3 Stage Paths

- * $RW \rightarrow OF$

- * 2 Stage Paths

- * $RW \rightarrow EX$

- * $MA \rightarrow OF$ (**X Not Required**)

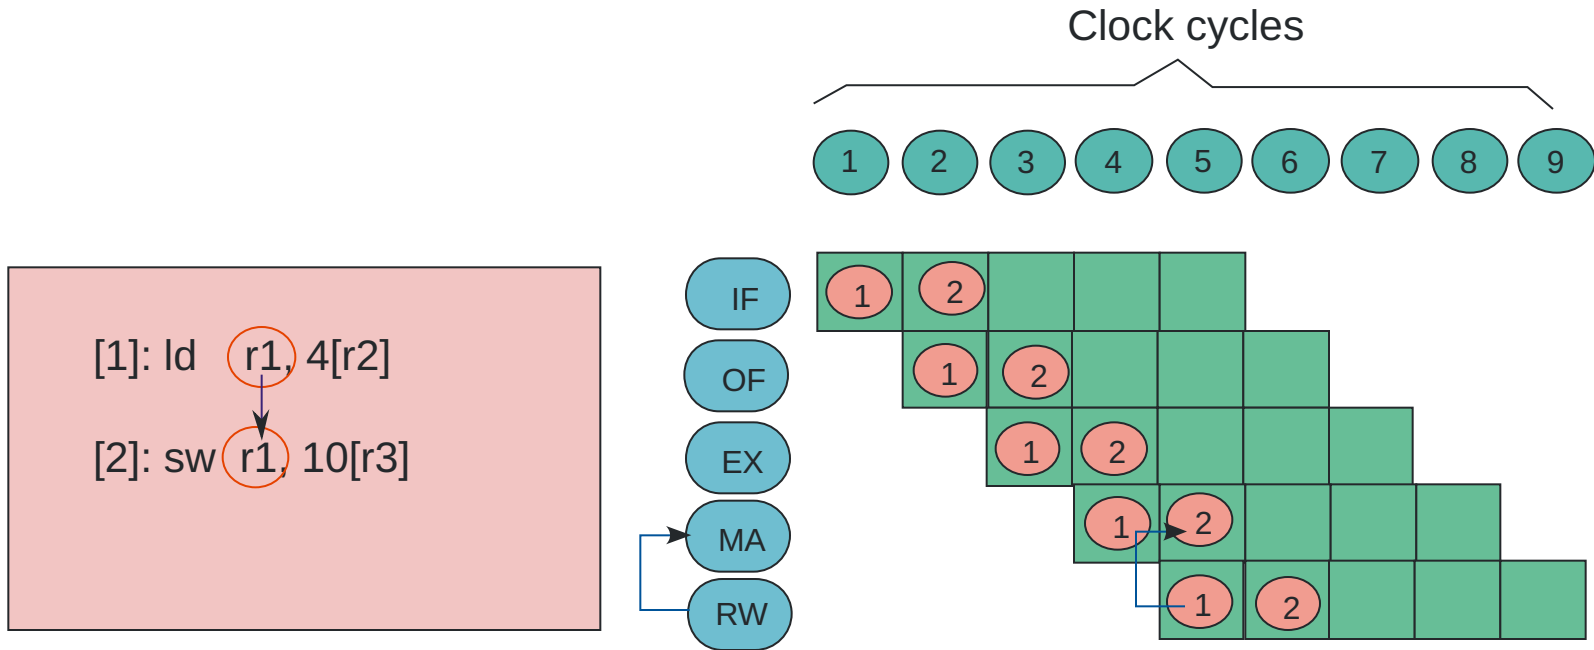
- * 1 Stage Paths

- * $RW \rightarrow MA$ (load to store)

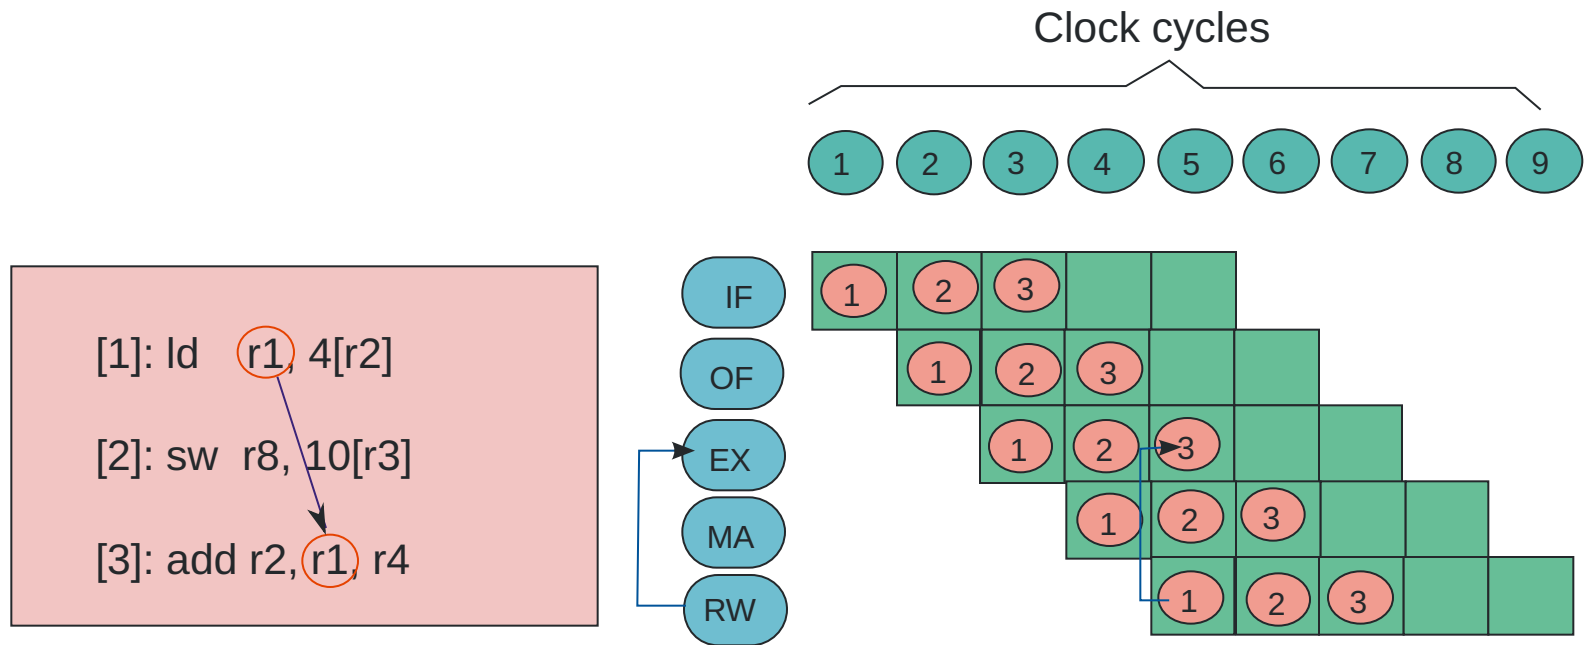
- * $MA \rightarrow EX$ (ALU Instructions, load, store)

- * $EX \rightarrow OF$ (**X Not Required**)

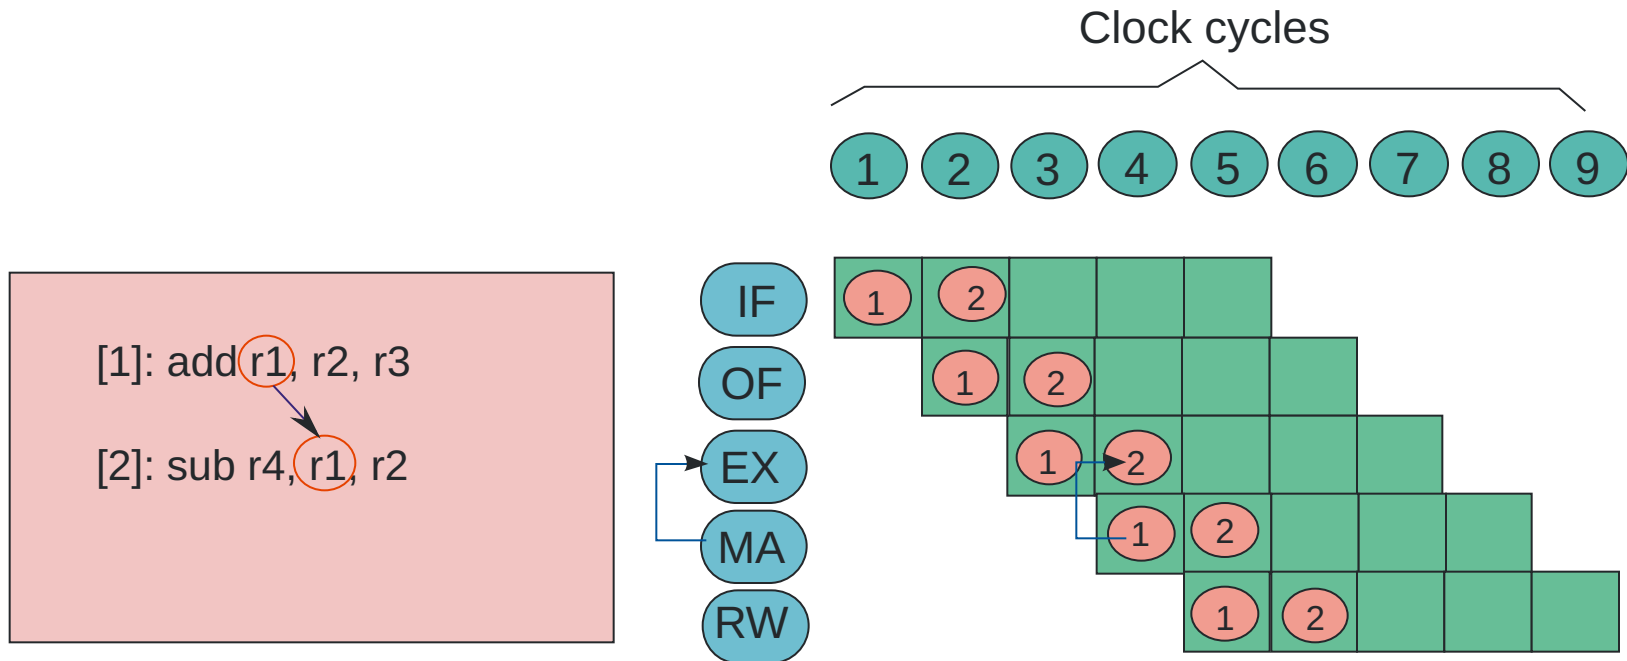
Forwarding Paths : RW → MA



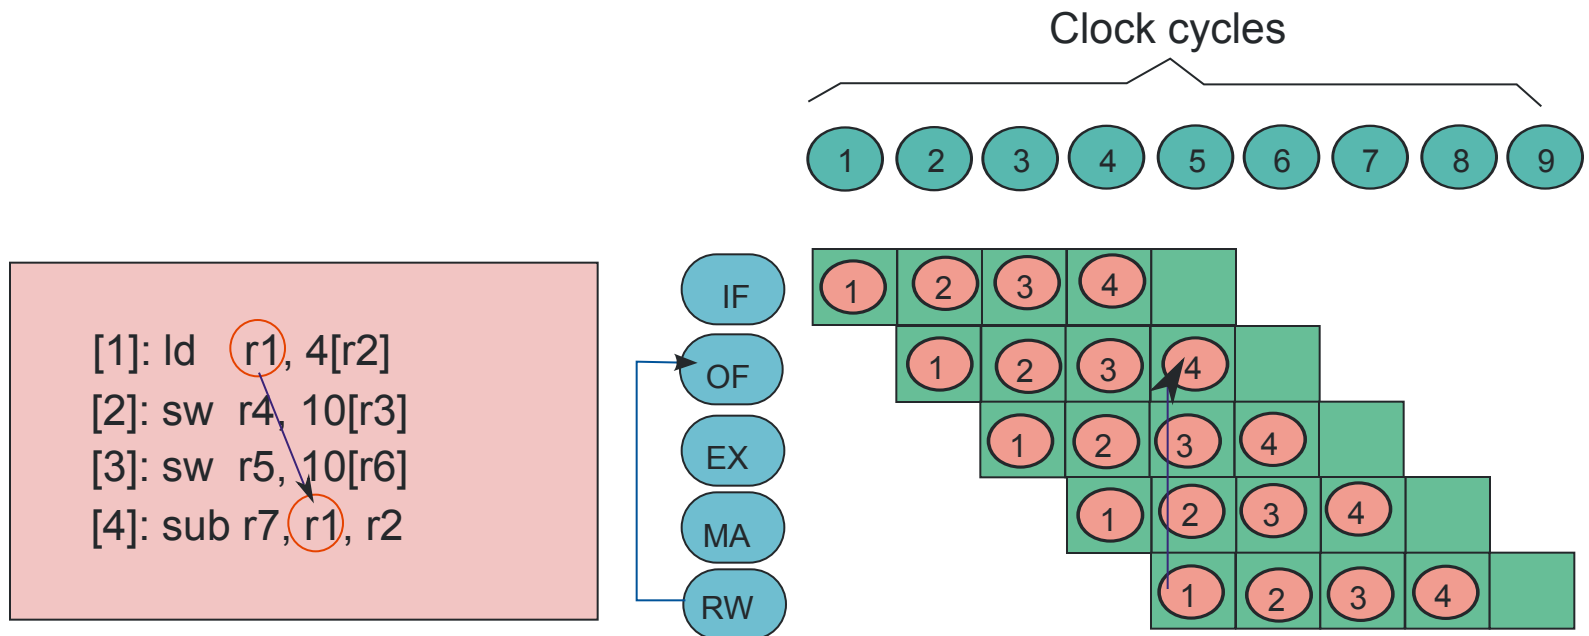
Forwarding Paths : RW \rightarrow EX



Forwarding Path : MA \rightarrow EX



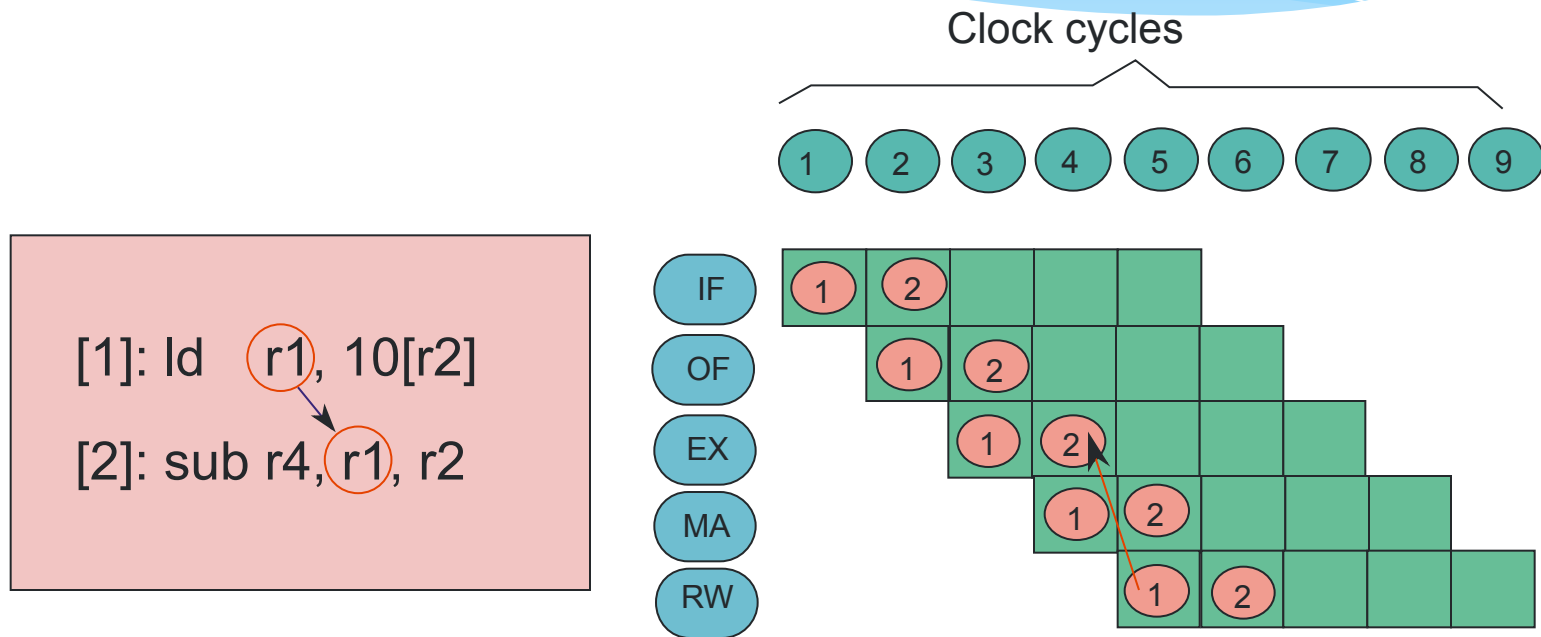
Forwarding Path : RW → OF



Data Hazards with Forwarding

- * Forwarding has unfortunately not eliminated all data hazards
- * We are left with one special case.
- * Load-use hazard
 - * The instruction immediately after a load instruction has a RAW dependence with it.

Load-Use Hazard

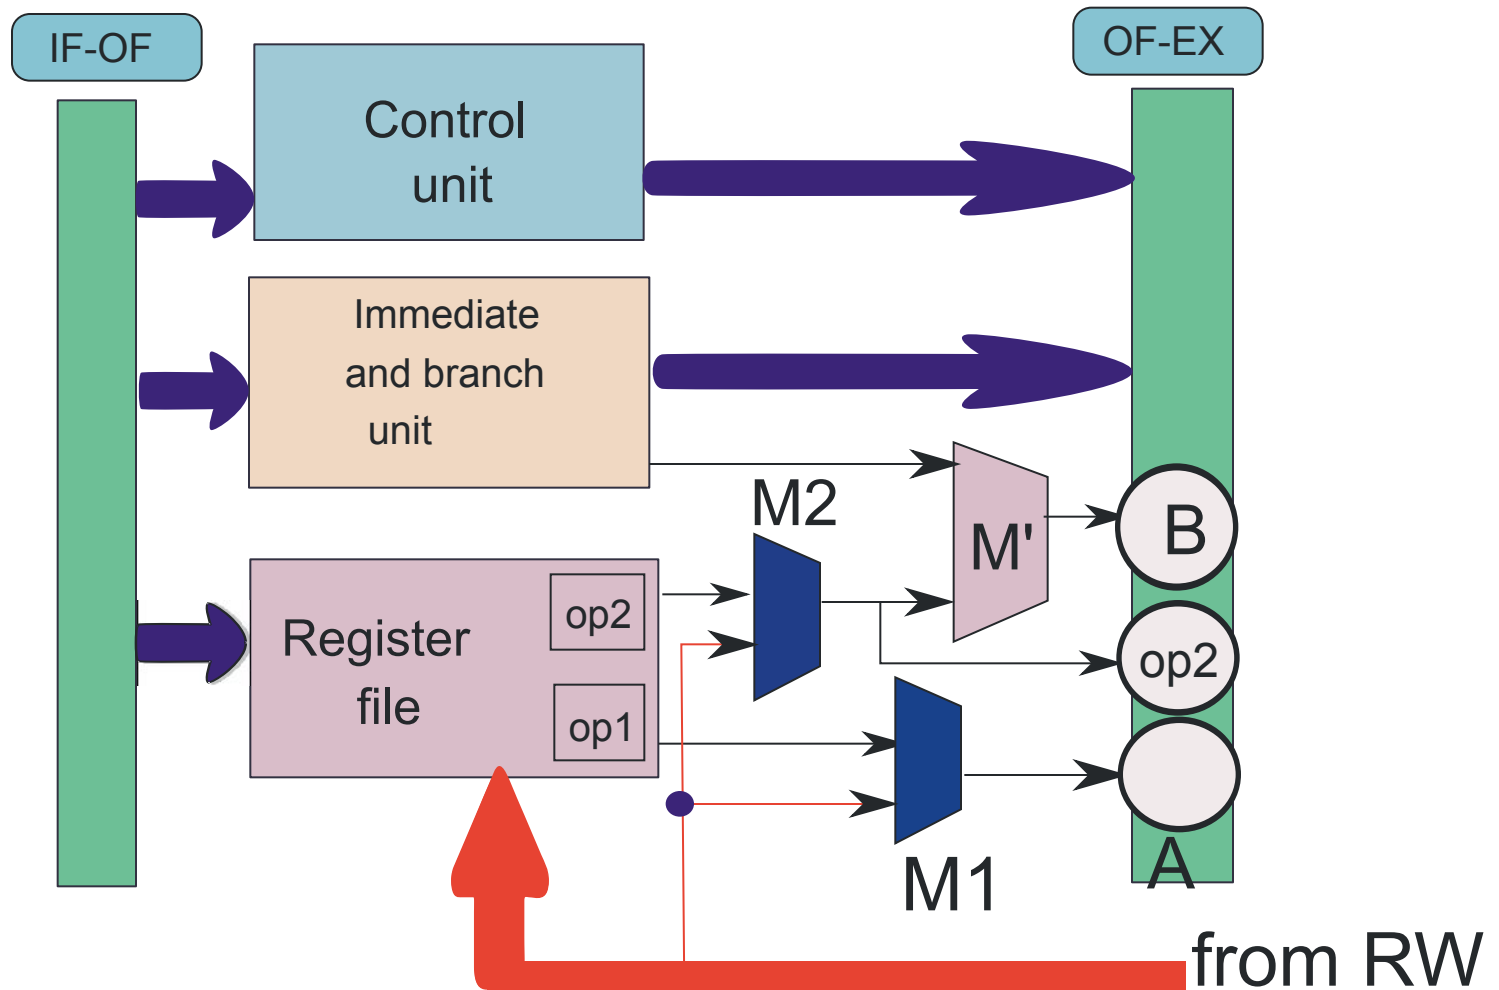


- * **Cannot forward** (arrow goes backwards in time)
- * Need to add a bubble (then use RW → EX forwarding)

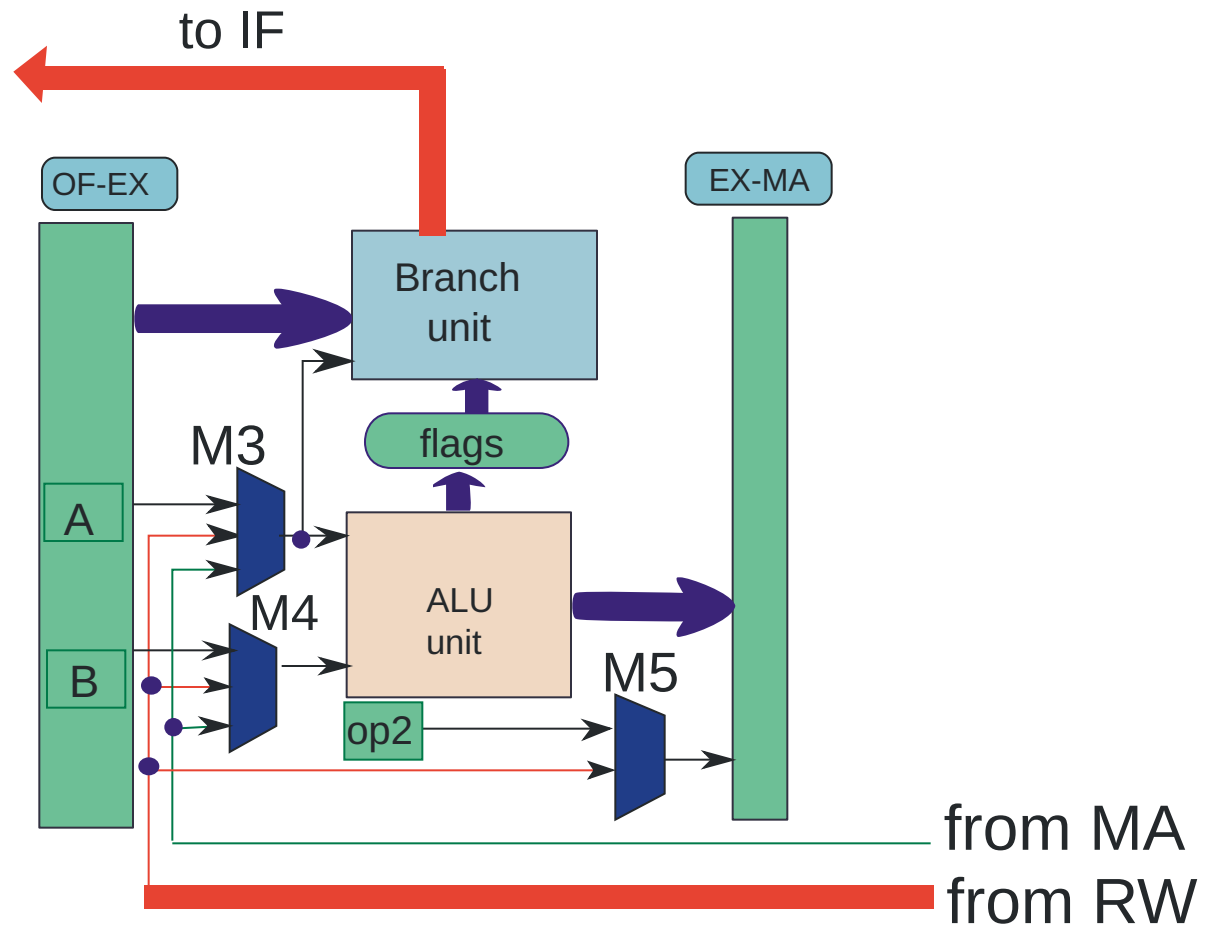
Implementation of Forwarding

- * At every stage there is a **choice** of inputs for each **functional unit**
 - * Use the **default** inputs from the previous stage
 - * **OR**, use one of the **forwarded** inputs
 - * Use a **multiplexer** to choose between the inputs
 - * A dedicated **forwarding unit** generates the signals for the **forwarding multiplexers**


OF Stage with Forwarding




EX Stage with Forwarding



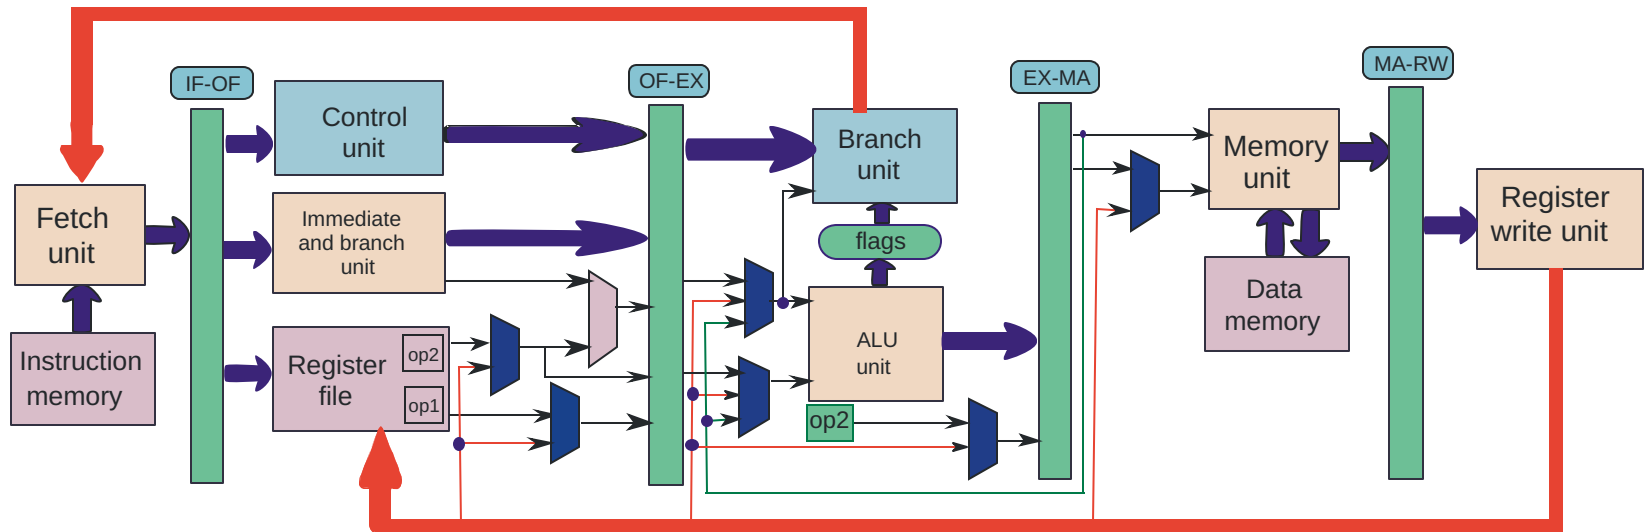
MA Stage with Forwarding



RW Stage with Forwarding



Data Path with Forwarding



Forwarding Conditions

- * Determine if there is a **conflict** between two **instructions** in different **stages**
 - * Find if there is a **conflict** for the first operand (rs1/ ra)
 - * Find if there is a **conflict** for the second operand (rs2/rd)
- * Always **forward** from the **latest** instruction (that is earlier than the current instruction)

Algorithm 6: Conflict on the first operand (rs1/ra)

Data: instructions, [A], and [B] (possible forwarding: [B] → [A])

Result: conflict exists on rs1/ra (**true**), no conflict (**false**)

if [A].opcode ∈ (nop, b, beq, bgt, call, not, mov) **then**

 /* Does not read from any register */

return false

end

if [B].opcode ∈ (nop, cmp, st, b, beq, bgt, ret) **then**

 /* Does not write to any register */

return false

end

/* Set the sources */

src1 ← [A].rs1

if [A].opcode = ret **then**

 src1 ← ra

end

/* Set the destination */

dest ← [B].rd

if [B].opcode = call **then**

 dest ← ra

end

/* Detect conflicts */

if src1 = dest **then**

return true

end

return false

Algorithm 7: Conflict on the second operand (rs2/rd)

Data: instructions, [A], and [B] (possible forwarding: $[B] \rightarrow [A]$)

Result: conflict exists on second operand (rs2/rd) (**true**), no conflict (**false**)

```
if [A].opcode  $\in$  (nop, b, beq, bgt, call) then
    /* Does not read from any register */
    return false
```

```
end
```

```
if [B].opcode  $\in$  (nop, cmp, st, b, beq, bgt, ret) then
    /* Does not write to any register */
    return false
```

```
end
```

```
/* Check the second operand to see if it is a register */
```

```
if [A].opcode  $\neq$  (st) then
    if [A].I = I then
        return false
    end
```


```
end
```

```
/* Set the sources */
```

```
src2  $\leftarrow$  [A].rs2
```

```
if [A].opcode = st then
    src2  $\leftarrow$  [A].rd
```

```
end
```



```
/* Set the destination */  
dest  $\leftarrow$  [B].rd  
if [B].opcode = call then  
    dest  $\leftarrow$  ra  
end  
/* Detect conflicts */  
if src2 = dest then  
    return true  
end  
return false
```

Interlocks with Forwarding


* Data-Lock

- * We need to only **check** for the **load-use hazard**
- * If the **instruction** in the EX stage is a **load**, and the **instruction** in the OF stage uses its **loaded value**, then **stall** for 1 cycle

* Branch-Lock

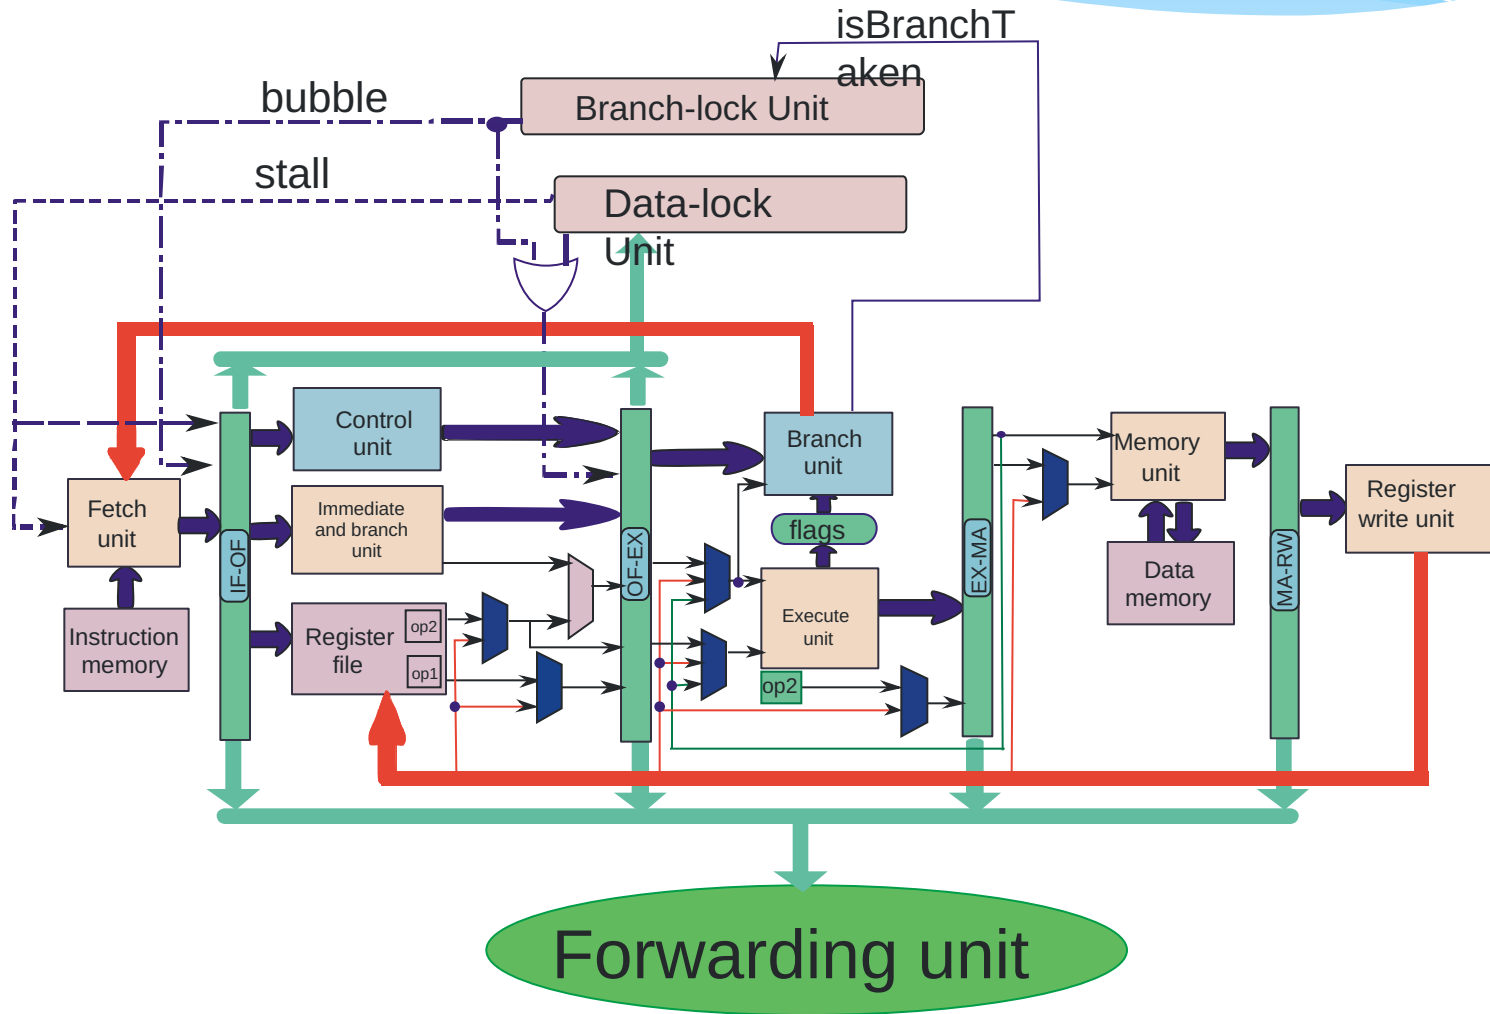
- * Remains the **same** as before.

The Curious Case of the *call* instruction




- * The **call instruction** generates the **value** of **ra** in the RW stage
- * Any instruction that uses the value of ra (such as ret) still works **correctly** !!!
- * Prove it

Complete Data Path



Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Measuring Performance

- * What do we mean by the **performance** of a processor ?
 - * **ANSWER** : Almost nothing
- * What should we ask instead ?
 - * What is the **performance** with respect to a given program or a set of programs ?
 - * **Performance** is **inversely proportional** to the time it takes to **execute** a **program**

Computing the Time a Program Takes

$$\begin{aligned}\tau &= \#seconds \\ &= \frac{\#seconds}{\#cycles} * \frac{\#cycles}{\#instructions} * (\#instructions) \\ &= \underbrace{\frac{\#seconds}{\#cycles}}_{1/f} + \underbrace{\frac{\#cycles}{\#instructions}}_{CPI} * (\#instructions) \\ &= \frac{CPI * \#insts}{f}\end{aligned}$$

- * CPI → Cycles per instruction
- * f → frequency (cycles per second)

The Performance Equation

$$P \propto \frac{IPC * f}{\#insts}$$

- * **IPC** → 1/CPI (Instructions per Cycle)
- * What are the **units** of performance ?
 - * **ANSWER** : arbitrary

Number of Instructions (#insts)

Static Instruction: The **binary** or executable of a **program**, contains a list of *static instructions*.

Dynamic Instruction: A **dynamic instruction** is a running instance of a static instruction, which is created by the **processor** when an instruction enters the pipeline.

- * Note that these are **dynamic** instructions
 - * **NOT** static instructions
- * A smart compiler can reduce the number of executed instructions

Number of Instructions(#insts) – 2

- * Dead code removal

- * Often **programmers** write **code** that does not determine the final output
- * This code is **redundant**
- * It can be identified and removed by the **compiler**

- * Function inlining

- * Very small **functions** have a lot of **overhead** → call, ret instructions, register spilling, and restoring
- * Paste the code of the **callee** in the code of the **caller** (known as **inlining**)

Computing the CPI

- CPI for a single cycle processor = 1
- CPI for an ideal pipeline (no hazards)
 - Assume we have n instructions, and k stages

Computing the CPI

- * CPI for a single cycle processor = 1
- * CPI for an ideal pipeline(no hazards)
 - * Assume we have n instructions, and k stages
 - * The first instruction enters the pipeline in cycle 1
 - * It leaves the pipeline in cycle k
 - * The rest of the $(n-1)$ instructions leave in the next $(n-1)$ consecutive cycles

$$CPI = \frac{n + k - 1}{n}$$

Computing the Maximum Frequency

- * Let the **maximum** amount of time that it takes to execute any **instruction** be :
 - * t_{\max} (also known as **algorithmic work**)
- * **Minimum** clock cycle time of a single cycle pipeline $\rightarrow t_{\max}$
- * In the case of a pipeline, let us assume that all the pipeline **stages** are **balanced**
- * Time per stage $\rightarrow t_{\max} / k$

Maximum Frequency - II

- * Let the **latch delay** be l
- * We thus have :

$$t_{stage} = \frac{t_{max}}{k} + l$$

$$\frac{1}{f} = \frac{t_{max}}{k} + l$$

The minimum cycle time ($1/f$) is equal to t_{stage} . Let us thus, assume that our cycle time is as low as possible.

Performance of an Ideal Pipeline

- * Let us assume that the number of instructions are a **constant**

$$\begin{aligned} P &= \frac{f}{CPI} \\ &= \frac{\frac{t_{max}}{k} + l}{\frac{n + k - 1}{n}} \\ &= \frac{\left(\frac{t_{max}}{k} + l\right) * (n + k - 1)}{n} \\ &= \frac{(n - 1)t_{max}}{k} + (t_{max} + ln - l) + lk \end{aligned}$$

Optimal Number of Pipeline Stages

$$\begin{aligned} \frac{\partial \left(\frac{(n-1)t_{max}}{k} + (t_{max} + l - l) + lk \right)}{\partial k} &= 0 \\ \Rightarrow -\frac{(n-1)t_{max}}{k^2} + l &= 0 \\ \Rightarrow k &= \sqrt{\frac{(n-1)t_{max}}{l}} \end{aligned}$$

- * k is inversely proportional to \sqrt{l}
- * k is proportional to $\sqrt{t_{max}}$

Implications

- * As we **increase** the **latch delay**, we should have **less** pipeline stages
 - * We need to **minimise** the time wasted in accessing latches
- * As we increase the **amount** of **algorithmic work**, we require more pipeline stages for ideal **performance**
 - * More pipeline stages help **distribute** the work better, and increase the **overlap** across instructions

A Non-Ideal Pipeline

- * Our ideal CPI ($CPI_{ideal} = 1$) is 1
- * However, in reality, we have stalls

$$CPI = CPI_{ideal} + stall_rate * stall_penalty$$

- * Let us assume that the **stall rate** is a function of the **program**, and its nature of **dependences**

Non-Ideal Pipeline - II

- * Let us assume that the **stall penalty** is **proportional** to the number of **pipeline stages**
- * Both these assumptions are strictly **not correct**. They are being used to make a coarse grained **mathematical model**.
- * $CPI = (n+k-1)/n + rck$
 - * $r \rightarrow$ **stall rate**, $c \rightarrow$ constant of **proportionality**

Mathematical Model

$$\begin{aligned}
 P &= \frac{f}{CPI} \\
 &= \frac{\frac{1}{\frac{t_{max}}{k} + l}}{\frac{n+k-1}{n} + rck} \\
 &= \frac{n}{\frac{(n-1)t_{max}}{k} + (rcnt_{max} + t_{max} + ln - l) + lk(1 + rcn)}
 \end{aligned}$$

Mathematical Model - II

$$\frac{\partial \left(\frac{(n-1)t_{max}}{k} + (rcnt_{max} + t_{max} + \ln - l) + lk(1 + rcn) \right)}{\partial k} = 0$$

$$\Rightarrow - \frac{(n-1)t_{max}}{k^2} + l(1 + rcn) = 0$$

$$\Rightarrow k = \sqrt{\frac{(n-1)t_{max}}{l(1+rcn)}} \approx \sqrt{\frac{t_{max}}{lrc}} \quad (as \ n \rightarrow \infty)$$

Implications

- * For programs with a lot of **dependences** (high value of **r**) → Use **less pipeline** stages
- * For a pipeline with **forwarding** → **c** is **smaller** (than a **pipeline** that just has interlocks)
 - * It requires a larger number of **pipeline stages** for optimal performance

Implications

- * The optimal number of pipeline stages is directly proportional to $\sqrt{(t_{\max} / I)}$
 - * This ratio is not significantly changing across technologies.
 - * This explains why the number of pipeline stages has remained more or less constant for the last 5-10 years

Example

Example Consider two programs that have the following characteristics.

Program 1		Program 2	
Instruction Type	Fraction	Instruction Type	Fraction
loads	0.4	loads	0.3
Branches	0.2	Branches	0.1
ratio(taken branches)	0.5	ratio(taken branches)	0.4

The ideal CPI is 1 for both the programs. Let 50% of the load instructions suffer from a load use hazard. Assume that the frequency of P_1 is 1, and the frequency of P_2 is 1.5. Here, the units of the frequency are not relevant. Compare the performance of P_1 and P_2 .

Example

Answer:

$$\begin{aligned} CPI_{new} = & CPI_{ideal} + 0.5 \times (\text{ratio}(\text{loads})) \times 1 \\ & + \text{ratio}(\text{branches}) \times \text{ratio}(\text{taken branches}) \times 2 \end{aligned} \quad (9.13)$$


We thus have:

$$CPI_{P_1} = 1 + 0.5 \times 0.4 + 0.2 \times 0.5 \times 2 = 1 + 0.2 + 0.2 = 1.4$$

$$CPI_{P_2} = 1 + 0.5 \times 0.3 + 0.1 \times 0.4 \times 2 = 1 + 0.15 + 0.08 = 1.23$$


The performance of P_1 can be expressed as $f/CPI = 1 / 1.4 = 0.71$ (arbitrary units). Similarly, the performance of P_2 is equal to $f/CPI = 1.5/1.23 = 1.22$ (arbitrary units). Hence, P_2 is faster than P_1 . We shall often use the term, arbitrary units, a.u., when the choice of units is irrelevant.

Performance, Architecture, Compiler



Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



What happens when you press a key ?

- * The **keyboard** logs the key press
 - * **Converts** the key to **ASCII** or Unicode
 - * **Sends** the code to the **processor**
 - * The **processor** thus receives an **interrupt**
 - * It **suspends** the current **program**
 - * Jumps to the **interrupt handler**.
 - * The **interrupt handler** draws the shape associated with the key
 - * The **processor** returns to execute the original **program**

Exceptions

- * **Exceptions** are generated when
 - * A **program** accesses an illegal address
 - * We try to **divide** 5/0
 - * We issue an **invalid instruction**
 - * ...
- * **Exception** are treated the same way as **interrupts**
 - * Jump to the **exception** handler
 - * Come back and start executing **programs**

Precise Exceptions

* Informal definition

- * We need to **return** to the original **program** at exactly the same **point**, at which we had left it
- * The execution of the **interrupt handler** should not disrupt the execution of the original **program** in any way. The outcome of the original program should be independent of the **interrupt** (unless the program caused an **exception**).

Precise Exceptions - II

- * **Formal Definition**

- * Let us number the **dynamic instructions** in a program :
 $I_1 \dots I_n$
- * Let us assume that an instruction **completes** after it either updates memory, writes to registers, or reaches the MA stage (cmp, b, beq, bgt, ret)
- * Let the last **program instruction** that **completes** before the **first** instruction in the **interrupt handler completes**, be I_k
- * Let all the **program** instructions that **complete** before the first instruction in the interrupt handler **completes**, be **C**

Precise Exceptions - III

$$I_j \in C \Leftrightarrow (j \leq k)$$

- * We need to ensure this **condition**.
- * Also no **instruction** of the form $I_{k'}$ ($k' > k$) should **complete** before all the instructions in the **interrupt handler complete**
- * After returning, we can seamlessly **execute** I_k (same instruction) or I_{k+1} (next instruction)

Marking Instructions

- * When an **interrupt** arrives, let us **mark** the instruction in the **MA** stage
- * Otherwise, if there is an exception
 - * We mark the instruction, as soon as it encounters a fault/exception
- * Once, an instruction is **marked**, we have two kinds of **instructions** in the pipeline
 - * Instructions **before/after** the **marked** instruction

Implementing Precise Exceptions

- * Wait till a **marked instruction** reaches the end of the **pipeline**
 - * Convert all the **instructions** after the marked instruction to **bubbles**
- * Ensures that both the **conditions** of a **precise exception** are met.
- * Once the **marked** instruction reaches the end of the **pipeline**, the exception unit loads the **pc** of the **interrupt handler**

Saving/Restoring Program State

- * Program State

- * PC
- * Registers
- * Flags
- * Memory

- * **Memory** → Assume that there is no **overlap** of memory regions, unless explicitly intended

oldPC Register

- * Let us add a **npc** field in the instruction packet
 - * For taken branches it is equal to the branch target
 - * For all other instructions it is equal to $(pc + 4)$
- * We populate the **npc** field in the EX stage
- * Depending on the type of the exception, we might want to return to **pc** or **npc**
 - * The **exception unit** sets the **oldPC** register to the right **return address**

Spilling/ Restoring Registers

- * In the case of **functions**, we stored registers on the **stack**
 - * In this case, the **interrupt handler** has a separate stack.
 - * We cannot **overwrite** the stack pointer (we will lose its previous value)
- * **Solution :**
 - * Use an additional register, **oldSP**, to save the stack pointer of the program.
 - * Load the new stack pointer, and spill all the registers
 - * Save oldPC

The Strange Case of the Flags

* Naive solution :

- * Do not allow any instruction after the marked instruction to update the flags register
- * We detect an exception typically towards the middle or end of a cycle
- * By that time, the instruction might have already updated the flags register (at least the master latch)

Solution

- * Add a **flags** field to the **instruction packet**
- * Every **instruction** saves the updated value of the **flags register** to the **flags** field in its instruction packet.
- * The exception unit saves the flags field of the marked instruction to the **oldFlags** register
- * We store the **oldFlags** register to a location in the interrupt handler's **stack**

Privileged Instructions

- * We add the following new **registers** to the ISA
 - * **special registers** → flags, oldFlags, oldPC, oldSP
 - * The **flags** register is now accessible to the ISA
- * Add the **z series** of privileged instructions
- * **movz** instruction
 - * Transfers values between **regular** registers and **special** registers

Privileged Instructions - II

- * **retz**

- * $pc \leftarrow \text{oldPC}$

- * The **z series** instructions can only be used by the interrupt handler, and the operating system

- * We use a **CPL** bit (current privilege level) bit for setting **permissions**

- * User programs (CPL = 1)

- * Interrupt handlers, kernel programs, (CPL = 0)

Implementing *movz* and *retz*

- * Define two new opcodes
 - * These *opcodes* are usable only when $CPL = 0$
 - * Augment the OF and RW stages to use the *special registers*.
 - * *movz*, and *retz* see a different view of registers.

Register	Encoding
<i>r0</i>	0000
<i>oldPC</i>	0001
<i>oldSP</i>	0010
<i>flags</i>	0011
<i>oldFlags</i>	0100
<i>sp</i>	1110

Assembly Code for Spilling Registers

```
/* save the stack pointer */
movz oldSP, sp
mov sp, 0x FF FC

/* spill all the registers other
than sp*/
st r0, -4[sp]
st r1, -8[sp]
st r2, -12[sp]
st r3, -16[sp]
st r4, -20[sp]
st r5, -24[sp]
st r6, -28[sp]
st r7, -32[sp]
st r8, -36[sp]
st r9, -40[sp]
st r10, -44[sp]
st r11, -48[sp]
st r12, -52[sp]
```


Spilling Registers - II

```
st r13, -56[sp]
st r15, -60[sp]

/* save the stack pointer */
movz r0, oldSP
st r0, -64[sp]

/* save the flags register */
movz r0, oldFlags
st r0, -68[sp]

/* save the oldPC */
movz r0, oldPC
st r0, -72[sp]

/* update the stack pointer */
sub sp, sp, 72
/* code of the interrupt handler */
....
....
....
```

Restoring Registers

```
/* update the stack pointer */
add sp, sp, 72

/* restore the oldPC register */
ld r0, -72[sp]
movz oldPC, r0

/* restore the flags register */
ld r0, -68[sp]
movz flags, r0

/* restore all the registers other than sp*/
ld r0, -4[sp]
ld r1, -8[sp]
ld r2, -12[sp]
ld r3, -16[sp]
ld r4, -20[sp]
ld r5, -24[sp]
ld r6, -28[sp]
ld r7, -32[sp]
ld r8, -36[sp]
```

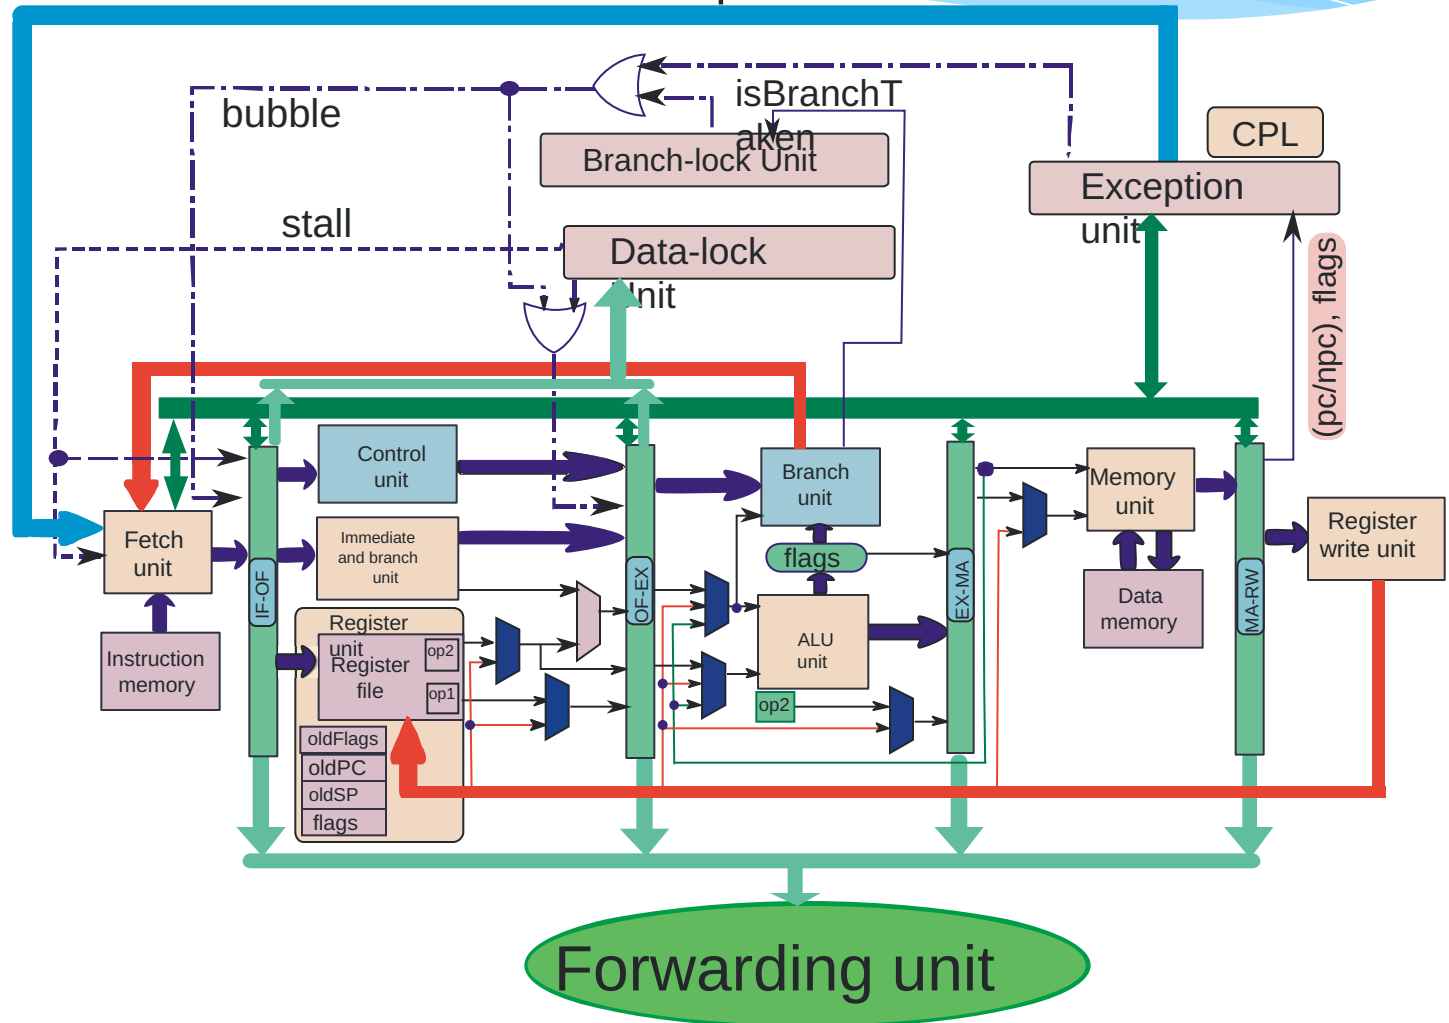
Restoring Registers – II

```
ld r9, -40[sp]
ld r10, -44[sp]
ld r11, -48[sp]
ld r12, -52[sp]
ld r13, -56[sp]
ld r15, -60[sp]

/* restore the stack pointer */
ld sp, -64[sp]

/* return to the program */
retz
```

PC of exception handler





THE END